

R Programming for Data Science

STSCI3040 - STSCI 5040

Fall 2023

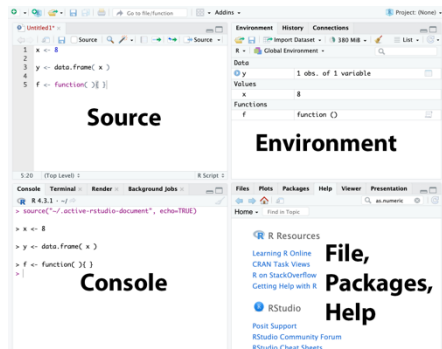
This page intentionally blank

Statistics courses usually use clean and well-behaved data, this leaves many unprepared for the messiness and chaos of data in the real world. This course aims to prepare students for dealing with data using the R programming language. The introduction will overview the basic R syntax, foundational R programming concepts such as data types, vectors arithmetic, and indexing, importing data into R from different file formats. The data wrangling topics include how to tidy data using the tidyverse to better facilitate analysis, string processing with regular expressions and with dates and times as file formats, web scraping, and text mining. Data visualization topics will cover visualization principles, the use of ggplot2 to create custom plots, and how to communicate data-driven findings.

Chapter 1: RStudio & R

To make use of **R** for data science, you need to install the **R** programming language on your computer. Versions of **R** can be found at <https://cran.r-project.org>. The installation comes with a program that acts as an **R**-language terminal. You type in a command and the command is executed. This program is all you need to write code in **R**. However, can be cumbersome to use. This is where **RStudio** come in. **RStudio** builds an environment around the basic **R** application. This environment consists of a series of panes that allow all the functionality of the **R**-programming language to be managed easier.

When using, **RStudio** will display four panes of information¹ These four panes are laid out as follows² :



Each of these panes can have several tabs. Not all of these tabs will be discussed.

Console

The **Console** pane will contain several different tabs. The most important of which is the **Console** tab. This tab acts as the stand-alone **R**-language terminal that was mentioned above. Essentially, it acts as a calculator. At the prompt, a command is typed in, once the enter/return key is pressed, the command is executed. This has the effect of giving you a very direct connection to executing **R** code, but it is not very efficient. One command needs to be entered at a time³. However, if an error is made, it may be the case that each preceding line needs to be reentered. This is what make the **Source** pane very useful.

Source

The **Source** pane will hold as many tabs as you can create. For the most part, each tab will either be a text editor or a table viewer. The table viewer shows the content of a dataset that has been loaded. The text editors will be the most useful. A new tab (text editor) can be opened in the **Source** pane by selecting *File* and then *New File* in the drop-down menus. Under *New File*, a large selection of options are available.

The option that is chosen under *New File* depends upon the purpose of the work produced. Each of the options will produce a text with with a different file extension when saved⁴. The extension that is attached to these files has an effect on the options that are visible in **RStudio**, how the contents are displayed, and the structure of how the content is written.

The two options under *New File* that will be focused on here:

- R Script
- R Markdown...

R Script:

An **R Script** is simple a text document with different **R** commands written on different lines⁵. As was said before, using the **Console** directly, each command must be typed in one at a time. Which means that each time a collection of commands needs to be executed, they need to be retyped into the **Console** prompt. With code written into an **R Script**, a *Run* command can be used to send all, or a selected portion, of an **R Script** into the **Console**. This alleviated retyping a collection of commands into the Console. The **R Script** can be saved as a text file with a *.R* extension. This allows to the set of commands to be saved and possibly transported to other locations. It should be kept in mind, that the output of an

¹ There will be times you only see three panes. Usually, this the first time you open **RStudio** or after you have closed all tabs in the **Source** pane.

² The layout of these panes can be changed in the *Global Options* of **RStudio** (under the *Tools* dropdown menu).

³ Multiples can be entered if they are separated by a semicolon

⁴ Even though these will not have a *.txt* extension, such files will be text files. They can be opened and edited in any text editor

⁵ Choosing **R Script** indicates that **R** code will be used. Selecting one of the other script options would indicate that a different programming language is being used.

executed **R Script** does not become part of the **R Script**. Its output will appear in the **Console** or one of the other panes in **R**.

R Markdown...:

Selecting **R Markdown...** will produce a text document with a *.Rmd* extension. Unlike an **R Script**, creating an **R Markdown** file is only a portion of the created output. With an **R Script**, the **.R* file is the output to be shared. An **R Markdown** file is only part of the output to be shared. An **R Markdown** file when properly processed can be used to produce a formatted Word document, a PDF, an HTML file for a report. Within an **R Markdown**, the plain text of a report can be written and using a special syntax **R** code can be included to perform computations and create graphics that are needed and incorporated into a report. All of this can be done with one document. There is no need to use one program to perform computations, a second program to make graphics, and a third to mix the two together to make a report. An **R Markdown** can contain all of these at once. Except for the datafiles, an **R Markdown** can contain all the information to create a report. The advantage of this is two fold:

- A reviewer only interested in the results can examine a version of the report (say a word document) with the code hidden (if desired).
- If a modification to the computations, graphics, or data is needed, the **.Rmd* file can easily be changed or data updated, and the **.Rmd* can be reprocessed to update the final report.

Environment

Before discussing the **Environment** pane, the idea of an **environment** needs to be explained. When using *R*, objects that are not built-into the system can be created. These objects may contain data, functions (collections of commands), or other information. These objects when created live in an **Environment** until that **environment** until they are removed from the environment or the environment is deleted. The interesting part of this, is that more than one **environment** can exist. The **Global Environment** is the most visible **environment**. Its contents are visible in the **Environment** tab in the **Environment** pane. The contents of the **Global Environment** are objects that are created when using the **Console** directly, running **R** code from an **R Script**, or directly running code within an **R Markdown** without processing the entire **R Markdown** to create a report⁶.

File/Packages/Help

This last pane has many useful tabs:

1. **Help**: This tab gives searchable access to the documentation for all the functions that are available for use in your code that you have not specifically written. The documentation usually indicates what inputs and outputs are needed and give examples/vignettes of using the object.
2. **Plots**: This tab is used to display plots that you have created. Options are available for saving them using menu options. If viewing plots is important, care should be taken to not make this pane too small. Otherwise, errors will appear or odd graphics can appear.
3. **Files**: This tab gives a file browser for the computer being used.
4. **Packages**: When initially installed, the **R** programming language has a limited set of capabilities. Other **R** users have taken those basic capabilities and used them to create packages. A package is a collection of new commands that can be loaded, if desired, and made available for use. Additional packages are optional. If a new capability is not needed, there is no reason to make it available. Some packages extend the basic linear regression capabilities, others give more graphics options, some add capabilities for formatting text. The **packages** tab keeps track of which packages a user has downloaded, which packages are currently active, and give options for downloading new packages.⁷

⁶ As was said, other **environments** can be made. A common way second **environment** is temporarily created is when an **R Markdown** is processed to create a report as a Word document, a PDF, or an HTML file. This second **environment** will not be visible to the user as the report is being created. However, it is important to be aware that it exists. As the report is being created, errors may occur because certain objects may exist in the **Global Environment** and not in the second **environment**. Because of this, it is important when writing **R Markdown** files, that all object that are used within it, are coded as part of the **R Markdown** file, before they are used

⁷ It is important to keep in mind that some packages might overwrite existing commands. Depending upon you need, you may want to use the version from a package or you may want to use the existing version. This is why there is a capability to load and unload packages.

Chapter 2: Basic RMarkdown

An **R Markdown** file is text file that can be processed to produce a Word, a PDF, or an HTML document. To process one correctly several packages need to be installed and loaded. The most important are **rmarkdown** and **knitr**. In fact, the process of producing the finished Word, a PDF, or an HTML document is usually referred to as **KNITTING**.

An **R Markdown** file will consist of three types of text:

- YAML header,
- plain text, and
- code chunks.

A brief example can be seen on the next page.

YAML Header

The **YAML header** is the first thing found in most **R Markdown** documents. It will be located at the top of the document with its contents surrounded above and below by three dashes set on their own line. Between the dashed lines several options can be set. To do so, the options name is listed followed by a colon and a space. Then the value for that option is typed it. Common options are used are: title, subtitle, author, date, and output. When an **R Markdown** document is knitted together to produce a Word, a PDF, or an HTML document, the values for the first four options will appear on the first page of the resulting output. However, the exact appearance will differ based upon the output values starting with one of these word_document, pdf_document, html_document⁸. On lines after the output line, other options can be added.

Code Chunks

Code Chunks are sections of an **R Markdown** file that contain **R** code. They are sections of text surrounded by two separate lines that start with (at least) three tick marks. The initial set of three tick marks are followed immediately by a set of curly braces. The contents of the curly braces, separated by commas are:

- the programming language being used. The code contained in a chunk can be any of a number of different programming languages. Therefore, the language being used needs to be identified. As a course in *R*, the identifier is a lower case **r**⁹.
- the name of the code chunk. This is optional. It is useful if a code chunk will be used more than once within a document. The name can be referenced later without recreating the entire code chunk. However, if two code chunks are named, they must be given different names.
- evaluation options

Three code chunks appear in the following example. They are *setup*, *cars*, and *pressure*. The *setup* code chunk is usually the first chunk that appears after the YAML header. Code within it is usually used to set default options for all subsequent codes. This is done by using the `knitr::opts_chunk$set()` command. The other two code chunks contain code to summarize and graph some data¹⁰.

When the text of an R Markdown file is displayed in the source window, it will be highlighted. Additionally, in the upper right hand corner of each code chunk are some buttons. The gear gives you a menu of some options that can be set for the current chunk. The downward facing arrow will execute all previous code chunks. The rightward facing arrow will execute the current chunk¹¹.

⁸ Typing in this line is not all that necessary. A **Knit** button appears at the top of the **Source** window when an **R Markdown** file is loaded. It has an option for the type of knitted output and it will fill it in automatically.

⁹ Different coding languages can not be mixed inside a code chunk.

¹⁰ The meaning of the code contained in these chunks will become apparent as the course progresses.

¹¹ When a R Markdown file is processed, code chunks are executed from the top down. Without knitting your code, you can use the first arrow to reload previous chunks, if they have been changed, while you work on the current chunk.

Text

The remainder of the example is the text that would make up a report. Careful examination will make apparent that this is not simply text. There are lines that start with pound signs (#) and words surrounded by symbols like asterisks or angles.(*,<,>). Usa of these symbols, and others, using the proper formatting will introduce headers, font formats, tables, and other items in the knitted result. In this example ****Knit**** indicates that the word Knit should be bold in the final result.

Example R Markdown Text

This text listed below is a basic template that appears when you create a new R Markdown file.

```
---
title: "Untitled"

author: "Jeremy Entner"

date: "8/25/2021"

output: word_document
```

```
---
```{r setup, include=FALSE}
knitr::opts_chunk$set(echo = TRUE)
```
```

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
```{r cars}
summary(cars)
```
```

Including Plots

You can also embed plots, for example:

```
```{r pressure, echo=FALSE}
plot(pressure)
```
```

Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

Chapter 3: Numerical Data

Numerical Data comes in THREE forms:

1. Integer
2. Floating Point/Double
3. Complex

We won't worry about complex numbers. An integer is any whole number, positive or negative. As an integer, it will be recognizable as a number followed by a capital L. For example, the integer three would appear as 3L.

A floating point/double number is any whole number or fractional number. The floating point/double version of the number three would simply appear as 3. The real difference between these two types is determined by how the computer stores the numbers. For the most part, we can ignore the difference between integers and floating point numbers.

Numerical Vectors

A **numerical vector** is a container in R that is used to hold numerical data. An individual numerical vector can contain zero or more numerical data values¹². Collections of numerical vectors can be combined using the **combine function** `c()`. This will produce a new vector whose elements are ordered as entered in `c()`. Make sure to separate each numerical data value with a comma.

```
c(1, 4, -10, 3.5)
[1] 1.0 4.0 -10.0 3.5
```

Mathematic Operations

Many standard mathematical operations can be performed on numerical vectors. The most basic would be addition.

```
1 + 2.3
[1] 3.3
```

While it looks like we added two number, what really happened was that R added together two numerical vectors that contained a single value each, namely 1 and 2.3.¹³ If we had wanted to add two longer vectors together, we could have done so

```
# Addition - ?Arithmetic
c(0, 2, 2.4) + c(1, -2, 0.6)
[1] 1 0 3
```

The plus sign `+` performs a **vectorized** addition. Corresponding pairs of values in each vector are added together. The first pair, 0 and 1, are added to give the first result, (1). The second pair, 2 and -2, are added to give the second result (0). Finally, the third pair, 2.4 and 0.6, result in (4).

Some basic mathematical operations that are vectorized are demonstrated below.¹⁴

```
# Subtraction - ?Arithmetic
c(0, 2, 2.4) - c(1, -2, 0.6)
[1] -1.0 4.0 1.8

# Multiplication - ?Arithmetic
c(0, 2, 2.4) * c(1, -2, 0.6)
[1] 0.00 -4.00 1.44

# Division - ?Arithmetic
c(0, 2, 2.4)/c(1, -2, 0.6)
[1] 0 -1 4
```

¹² at times each individual value may be called an **element**

¹³ The combine function is not needed for single values.

¹⁴ Multiple operations can be combined using parentheses. The order of operations can be described using PEMDAS: Parentheses/Grouping, Exponents, Multiplication, Division, Addition, Subtraction

Some more complex operations are also vectorized. However, for clarity, we will demonstrate them with single valued vectors.

```
# Absolute Value - ?MathFun
abs(-7.1)

[1] 7.1

# Square Root
sqrt(4)

[1] 2

# Exponents - ?Arithmetic
3^2

[1] 9

# Natural Log & Exponential- ?Log
log(4.8)

[1] 1.568616

exp(1.569)

[1] 4.801844

# Trig functions - ?Trig
cos(2)

[1] -0.4161468

# Quotient and Remainder - ?Arithmetic
10%/%3

[1] 3

10%%3

[1] 1
```

Recycling

Some operations with two inputs **recycle** (repeat) elements the shorter vector, until it is long enough to perform the vectorized operation.

```
c(0, 2, 3.4, 10) + c(1, 0)

[1] 1.0 2.0 4.4 10.0
```

In the first line, R repeated the vector with two elements. The second vector was treated like `c(1, 0, 1, 0)`.

```
c(0, 2, 3.4, 10) + c(1, 0, 7)

Warning in c(0, 2, 3.4, 10) + c(1, 0, 7): longer object length is not a multiple of shorter object length

[1] 1.0 2.0 10.4 11.0
```

In this example, a warning is given. It indicates that the longer vector is not a multiple of the shorter vector, in terms of length. Three does not divide evenly into four. The vector of length three was partially recycled as `c(1, 0, 7, 1)`.

Two Interesting Facts

1. The number of digits displayed can be controlled. However, It doesn't change the representd value
2. The number π is stored in R. It can be accessed by typing `pi`.

```
options(digits = 20) #?options
pi #?Constants

[1] 3.141592653589793116
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Using the help menu, determine the commands for computing the sine and tangent of a value.
2. Determine the following:
 - a. $19^{2.2}$
 - b. $\frac{19^{\ln 9}}{18} + 15$
 - c. $\sqrt{24}$
 - d. $\sin(6\pi)/\tan(7\pi)$
3. Which produce equal/identical results.
 - a. `c(12, 23, 46, 8-1)`
 - b. `c(1 π , 2 π , 3 π , 10 π)`
 - c. `c(1, 2, 4, 8)^c(2, 3, 6, -1)`
 - d. `c(12, 5, -4, 2) + c(1, 2)`
 - e. `c(10, 2, 3, 1)* π`
 - f. `c(12, 5, -4, 3) + c(1, 2, 1)`
4. Convert these temperatures (32, 212, 100, -40) from Fahrenheit to Celsius. Do this with one line of code. (You should look up the conversion formula.)
5. Compute the following using two vectors:
 - a. The quotient of 5 divided by 3. The quotient of 2 divided by 4. The quotient of 228 divided by 8.
 - b. Two squared. Three cubed. Four to the forth. Five fifth.
6. Compute the following making use of recycling.(Shrink the given vectors)
 - a. `c(1, 2, 3, 4, 5, 6) * c(0, 4, 6, 0, 4, 6) - c(1, 1, 1, 1, 1, 1)`
 - b. `c(1, 2, 3, 4, 5, 6) * c(0, 4, 6, 0, 4, 6) - c(1, 1, 1, 1, 1, 1, 1)`
 - c. `c(1, 2, 3, 4, 5, 6) * c(0, 4, 6, 1, 9, 0) - c(1, 1, 1, 1, 1, 1)`
7. Compute the following without using recycling.(Expand the given vectors as needed)
 - a. `c(1, 2) + c(0, 10, -1) * 20`
 - b. `c(1, 2) + (c(0, 10, -1) * 20)`
 - c. `c(1, 2) / c(1, 2, 1, 2) - c(1, -1, 1, -1, 1, -1)`

This page will eventually be filled with problems also.

Chapter 4: Assignment Operator

The **assignment operator** `<-` allows you to efficiently keep track of values used in computations and results produced by computations. Consider adding these vectors `c(0,2,2.4)` and `c(1,-2,0.6)`

```
c(0, 2, 2.4) + c(1, -2, 0.6)
[1] 1 0 3
```

The result is `c(1,0,3)`. Suppose you wanted to save this result or perform further computations using it. The operator will allow you to do this. In particular, it will allow this without the need to retype the vector.

`c(1,0,3)` is short, only three values. A result with 1000 values would not be very easy to retype.

The assignment operator can be used to attach a variable name to the result.

```
output <- c(0, 2, 2.4) + c(1, -2, 0.6)
```

To read this line of code, start on the right of the `<-` sign. First, `c(0,2,2.4) + c(1,-2,0.6)` is computed. Second, the assignment operator `<-` assigns the result to a variable called `output`. Notice two things about this:

1. The result is not displayed.
2. A new value called `output` appeared in the environment pane.



Once a variable name has been assigned to a value, the two become synonymous. When executed, the variable `output` produces the values stored in it.

```
output
[1] 1 0 3
```

Operations can be performed on it, functions can be applied to it, as if the original vector of values was being used.

```
output + 1 # The vector 1 is being recycled.
[1] 2 1 4
sin(output) # The sine function is vectorized.
[1] 0.841471 0.000000 0.141120
```

In fact, a variable can be used to update itself and add new values.

```
output
[1] 1 0 3
output <- c(output, 999) # a fourth element is added to output
output
[1] 1 0 3 999
```

The right hand side of `<-` takes the original values associated with the `output` variable, adds one to all of them, and then associates the result with the variable `output`. Effectively, the old values are erased, and replaced.



Several Variables and Multiple operation

Below, we assign variable names, perform some operations using the variables. The results have variable names assigned to them. Then, an operation is performed on the constructed results.

```
x <- c(0, 2, 2.4) # Assign the variable name x to this vector
y <- c(1, -2, 0.6) # Assign the variable name y to this vector
output.add <- x + y # Add x and y together, and assign a name to the result.
output.difference <- x - y # Subtract y from x, and assign a name to the result.
output.add * output.difference # Multiply output.add and output.difference together

[1] -1.0 0.0 5.4
```

(Theses operations were selected for demonstration purposes only.)

Rules for Variable Names

There are some rules for variable names that must be followed:

1. The name must start with a letter.
2. The name can only contain letters, numbers, underscores, and periods.

There are also some general guidelines for variables names. They will make reading and updating code much easier.

1. If a value, (vector of values, or any object that you can construct) is going to be used multiple times, assign it a variable name. If you discover a mistake in that value, correct the mistake where the variable name is assigned. No need to search all of your code.
2. The variable name should be assigned **before** every instance of it being used in your code. R reads code from top to bottom.

```
test.variable # This variable will not exist until the next line.

Error in eval(expr, envir, enclos): object 'test.variable' not found

test.variable <- 3
test.variable

[1] 3
```

3. Use a name that describes what the values represent. At a later time, the descriptive name makes it easier to understand the variables use.
4. Pick a convention for writing names and use is consistently.

```
variablename VARIABLENAME #all Lower/UPPER case, words not separated, hard to read
variable.name #all lower case, words separated by a period, easy to read
variable_name #all lower case, words separated by an underscore, easy to read
variableName #Camel case, words not separated, Capitalize first letter of each word except possible the first, easier to read
```

5. Select names for your variables that do not already have some meaning in R.

```
pi # This is a constant that is already defined in R.

[1] 3.141593

pi <- 10 # This assigns a new value to the name pi
pi

[1] 10
```

Other Assignment Operators

Other assignment operators exist.

1. The single equal sign =. This works like <-, but doesn't visually does not indicate what is occurring as well as <-.
2. -> assigns left-hand values into a variable listed on the right. If the value is long, it is hard to find the variable name.
3. <<-, ->> are assignment operators that **may** be used when making functions.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

- Create a single variable for each of the following vectors. Use a descriptive name. The naming convention to use is indicated.
 - 12 (period separator)
 - `c(1, 3, 5, 7, 9, 11, 13)` (All Capitals)
 - `c(1, 2, 3, 4, 5, 6, 7) * pi` (underscore separator)
 - `1 / c(1, 2, 2^2, 2^3, 2^4, 2^5, 2^6, 2^7, 2^8)` (Camel Case)
- Create the following variables, and assign an appropriate vector to it. Modify the variable name to use the indicated naming convention
 - `sqrt.ofTWO` (period separator)
 - `naturallogofonethroughten` (Camel Case)
 - `sine_of_0_45_90____360_degrees` (period separator)
 - `remaindersOfTen.DividedByZeroToTen` (All Capitals)
 - `cUBe_ROOTOf.3` (All Capitals)
 - `PO.WER_so.F5` (Camel Case)
 - `A_bs01u_TeV4alues` (period separator)
- Create a vector with variable name **AssignmentData001**. Use functions of **AssignmentData001** to compute the vectors that follow. Assign variable names to each result. Display each result. The elements of **AssignmentData001**, in order, should be:

1 4 7 9 12 15 19 21 22 30

- `c(1, 4, 7, 9, 12, 15, 19, 21, 22, 30) * pi`
- `1/c(3^1, 3^4, 3^7, 3^9, 3^12, 3^15, 3^19, 3^21, 3^22, 3^30)`
- `c(1, 1, 1, 1, 1, 1, 1, 1, 1, 1)`
- Given Newton's Law of Gravitation $F = G \frac{m_1 m_2}{r^2}$ where the gravitational constant $G \approx 12.6$ (In Oan Units), m_1, m_2 are the masses of two bodies, r is the distance between them, and F is the gravitational force between them, determine the gravitational force between bodies with the following properties:

| m_1 | m_2 | r |
|-------|-------|-----|
| 10 | 11 | 2 |
| 34 | 21 | .5 |
| 4 | 39 | 19 |
| 884 | 73 | 98 |
| 4 | 34 | 23 |

- Compute the first 20 numbers in the Fibonacci Sequence as a single vector. Store your result in a vector called `Fib`. Start with the number one.

Use the fact that the n^{th} value in the sequence F_n can be computed using

$$F_n = \frac{\psi^n - (-\psi)^{-n}}{\sqrt{5}}$$

where $\psi = \frac{1+\sqrt{5}}{2}$.

6. The following table provides the length of the hypotenuse and the length of another side of a right triangle. Determine the length of the third side for each of the following.

| Hypotenuse | Other |
|------------|-------|
| 5 | 4 |
| 13 | 12 |
| 41 | 40 |
| 29 | 21 |
| 65 | 63 |

7. Use the given variables to create the indicated vectors.

| VariableName | Vector |
|--------------|------------------|
| ones | c(1,1,1) |
| evens | c(2,4,6,8) |
| negatives | -c(10,20,30,120) |
| ninety_nine | 99 |

- c(1, 1, 1, 1, 1, 1, 99)
- c(1, 1, 1, 2, 4, 6, 8, 10, 20, 30, 120)
- c(1,1,1,99,10,20,30,120,10,20,30,120,1,1,1,1,1)
- c(2,2,2,4,8,12,16)
- c(0, 0, 0, 5, 5, 5, 15, 0, 0, 0, 0)

Chapter 5: Character Data

Character Data appears in two forms. It will appear as text surrounded by:

1. "Double quotes" or,
2. 'Single quotes'.

Commonly, character data are referred to as **string** data, or **strings**. Strings can contain any character you can type on the keyboard. Strings can even contain no characters. This is indicated by a set of quotes without any characters or spaces in between them.

```
" " #This is an empty string
[1] ""
" " #This string contains a space.
[1] " "
"This is a string. It contains words."
[1] "This is a string. It contains words."
```

Character Vectors

Collections of character data can be held in an ordered character (atomic) vector, in the same way as logical and numerical data. An ordered character vector can be created using the **concatenate function** `c()`. Make sure to surround each string with quotes and separate each string with a comma.

```
c("String", "Elements", "Are", "Wrapped", "In", "Quotes.", 9.7, 16)
[1] "String" "Elements" "Are" "Wrapped" "In" "Quotes." "9.7" "16"
```

Operating on Strings

The `paste()` function is used to join character vectors together element by element. The `sep =` argument indicates which string is used to separate the combined elements. `paste()` produces a character vector.

```
paste(c("This", "a", "vector"), c("became", "single", "."), sep = "<:>")
[1] "This<:>became" "a<:>single" "vector<:>."
```

Adding the `collapse =` argument will combine all elements into a single string.

```
paste(c("This", "a", "vector"), c("became", "single", "."), sep = "<:>", collapse = " - - ")
[1] "This<:>became - - a<:>single - - vector<:>."
```

Printing Strings

When entered on its own, a string will automatically display itself. However, in some situation, you must force the display. This can be done using several functions.

```
print("This is a string.")
[1] "This is a string."
noquote("This is a string without printed quotes.")
[1] This is a string without printed quotes.
```

The `cat()` function can be used to print a string without the line number printed and without quotes. It will also interpret formatting symbols.

```
cat("This line has a \t tab in it. \n Now, we have skipped a line and included è.")
This line has a      tab in it.
Now, we have skipped a line and included è.
```


Modifying Strings

Strings can be modified in several ways. The case of the letters in string can be changed. (These functions may not give the same results on two different computers, if the underlying locales are set differently.)

```
tolower("ChanGE cAsE 1234")
[1] "change case 1234"
toupper("ChanGE cAsE 1234")
[1] "CHANGE CASE 1234"
```

Character replacement can be performed with `chartr()`. The *old* argument gives the old values to be replaced. The *new* argument indicates what they are translated into. The final string is what need to be translated.

```
chartr(old = "EhC", new = "eHX", "ChanGE cAsE 1234")
[1] "XHanGe cAse 1234"
chartr(old = "abcX", new = "DEFx", "abcdefghijklmnpqrstuVwXyz")
[1] "DEFdefghijklmnpqrstuVwxyz"
```

Extracting Strings

It may occur that just a single string contains several pieces of information. You may want to separate these pieces or extract the single piece that you need.

To extract part of a string, use the `substr()` command. It requires that you know the numerical position of the characters that you want in the string (including spaces).

This code will extract the second, third, fourth and fifth value from the given string.

```
substr("abcdefghijklmnpqrstuVwXyz", first = 2, last = 5)
[1] "bcde"
```

In some cases, the portion that you want to extract might not have a known length. However, you may know that there will only be a certain number of characters, say 4, after it. Use the `nchar()` function to determine the number of characters in the string, and subtract 4 from this value.

```
string.with.equal.signs <- "=I Want This Part.=="
string.length <- nchar(string.with.equal.signs)
substr(text = string.with.equal.signs, first = 2, last = string.length - 3)
[1] "I Want This Part."
```

`substring()` can also be used to replace one part of one substring with another.

```
substring(text = string.with.equal.signs, first = 2, last = string.length - 3) <- "+++++"
string.with.equal.signs
[1] "=+++++ This Part.=="
```

If there is a regular pattern to how the information is given in the string, it can be split apart.

```
strsplit("Age , Height , Weight , Other", split = " , ")
[[1]]
[1] "Age" "Height" "Weight" "Other"
```

Two things we should notice about this:

1. The string was split using " , " instead of " , ". The first would leave a spaces on each resulting string.
2. The result has two sets of indexes `[[1]]` and `[1]`. This is because the result is not a vector, but a different object for holding data called a **list**.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Create a character vector called **blank**. It should have 5 elements. The middle three should be spaces. The outer two should be empty.
2. Create the following character vectors based on the given descriptions and assign it to a variable. Each individual bold item in the description should be its own element. The variable name should be descriptive of the content.
 - a. vowels (5 elements) (You may need to look up what a vowel is.)
 - b. The digits 0 through 9 (10 elements)
 - c. First five **words** in US National Anthem
 - d. First two **sentences** in A Tale of Two Cities.
3. Make a single string out of each of the following vectors.
 - a. `MoreCharacterData001`; link using spaces
 - b. `MoreCharacterData002`; link using underscores
4. Split this string `CharacterData003` into a character vector with one number per element (9inches becomes 9, 12inches becomes 12, etc). Leave no extraneous spaces or characters. This actually creates a list that contains one character vector, but that is ok.

5. Make a character vector with 26 elements. The first element should be "I like the letter A.". The second element should be "I like the letter B." This should continue until the twenty sixth element "I like the letter Z." Pay attention to the period. Make use of the vector CharacterData004.
6. Make a **single** string using the vectors CharacterData005 and CharacterData006. The final string should contain the ending of 'MacArthur's Farewell Address to Congress'. There should be no extraneous spaces at the end. It will start and end as follows:

I am closing my Good Bye.

7. A Caesar Cypher is a simple method for encoding a message. The method takes each letter in a given message and shifts it to a letter a fixed number of places to the right of it in the alphabet. Letters at the end of the alphabet wrap around to the beginning. For example, if the shift is 1, then A becomes B, B becomes C, Y becomes Z, and Z becomes A. Other shifts can be used. The messages below have been encoded with a Caesar Cipher. Decode them.
 - a. CharacterData007
 - b. CharacterData008
 - c. CharacterData009
 - d. CharacterData010
8. Extract the portion of each string that does not contain equal signs:
 - a. CharacterData011
 - b. CharacterData012
 - c. CharacterData013

Chapter 6: More Character Operations - stringr

The *stringr* package adds a number of functions for manipulating strings. Functions built into the *stringr* package are all named beginning with “str_”. Additionally, the arguments appear in a consistent order from function to function.

```
library(stringr)
```

Joining & Splitting Strings

The `str_c()` function joins multiple strings into one.

```
str_c("Number: ", c(1, 2, 3, 4, 5, 6))

[1] "Number: 1" "Number: 2" "Number: 3" "Number: 4" "Number: 5" "Number: 6"

str_c(c("a", "b", "c", "d", "e", "f"), c(1, 2, 3, 4, 5, 6), sep = "_")

[1] "a_1" "b_2" "c_3" "d_4" "e_5" "f_6"

str_c(c("a", "b", "c", "d", "e", "f"), c(1, 2, 3, 4, 5, 6), "X", sep = "_", collapse = "^")

[1] "a_1_X^b_2_X^c_3_X^d_4_X^e_5_X^f_6_X"
```

The `str_c()` function works in much the same way as `paste()`.¹⁵ `sep` is used to insert a string between the elements. `collapse` is used to combine all resulting elements into a single string.

The `str_flatten()` function will collapse the separate elements of a vector into a single string.

```
str_flatten(string = c("pear", "orange", "berry", "nut"))

[1] "pearorangeberrynut"

str_flatten(string = c("pear", "orange", "berry", "nut"), collapse = " - ")

[1] "pear - orange - berry - nut"
```

The `collapse` argument indicates the string to be inserted between each element.

The `str_split()` function breaks a string into several substrings based upon some given pattern.

```
str_split(string = "pear - orange - berry - nut", pattern = " - ", n = 3, simplify = FALSE)

[[1]]
[1] "pear"      "orange"    "berry - nut"
```

The `str_dup()` function is a vectorized function duplicates and concatenates strings. Its second argument *times* indicates how many times each element of the given vector is duplicated.

```
str_dup(string = c("pear", "orange", "berry", "nut"), times = c(1, 2, 3, 4))

[1] "pear"      "orangeorange"    "berryberryberry" "nutnutnutnut"
```

The `str_length()` function will return the number of characters in each element of a given vector. This is similar in function to `nchar()`.

```
str_length(string = c("pear", "orange", "berry", "nut"))

[1] 4 6 5 3
```

Substrings - Extracting, Modifying & Padding

In the *stringr* package, the `str_sub()` function extracts substrings from a character vector. The substrings will start with the character in the position indicated by *start* and end with the character in the position indicated by *end*. Failure to indicate either will be seen as an indication to select from the first character, or continue to the final character.

```
x <- c("abcdefg", "ABCDEFGH", "123456789")
str_sub(string = x, start = 2, end = 5)
```

¹⁵ With the exception that it will only recycle vectors of with only a single element. Otherwise, vectors must have the same length. Notice the error in the second example.

```
[1] "bcde" "BCDE" "2345"

str_sub(string = x, end = -4)

[1] "abcd" "ABCDE" "123456"

str_sub(string = x, start = 2, end = -2) <- "A"
x

[1] "aAg" "AAH" "1A9"
```

A negative number in the *end* position will indicate that the selection should be made by counting backwards from the end of a given string.

The *str_pad()* function can be used to add characters to a string until the string attains a minimum length. The example below will add some characters ("Z") to both sides of these foods until they have a minimum length of 6 characters.

```
str_pad(string = c("pear", "orange", "berry", "nut"), width = 6, side = "both", pad = "Z")

[1] "ZpearZ" "orange" "berryZ" "ZnutZZ"
```

Characters can be added solely to the “left” or “right” sides of the given strings. The function is vectorized in the *pad* argument, which lists a vector of single characters used to extend the given strings.

Whitespace - Trimming & Squishing

When strings are split into substrings, extraneous white spaces are included in the included substrings.

The *str_trim()* function will remove extraneous white spaces on the “left” side, “right” side, or “both” sides of the given string depending upon the argument *side*.

```
x <- " <--Extra Spaces Here--> <--Here--> "
str_trim(string = x, side = "both")

[1] "<--Extra Spaces Here--> <--Here-->"

str_squish(string = x)

[1] "<--Extra Spaces Here--> <--Here-->"
```

The *str_squish()* function removes all extra white spaces on either end of, and inside of, a given string. In the examples, pay close attention to the placement of the quotation marks.

Truncation

The *str_trunc()* function will truncate a given character string, and pastes an ellipsis onto the result. The *width* argument determines the length of the overall result. The *width* minus the length of the *ellipsis* argument determines the number of characters kept from the original string. The *ellipsis* can be placed on the “right”, “left”, or “center”.

```
x <- "This string has a lot of characters in it."
str_trunc(string = x, width = 11, side = "right", ellipsis = ".....")

[1] "This s....."
```

Sorting

The *str_sort()* function sorts a character vector. The *decreasing* argument determines whether the sorting is in decreasing or increasing order. *locale* defaults the sort order to English alphabet. This will affect the ordering used. *na.last* indicates where NA values fall in the ordering. *numeric* controls whether digits are treated as numbers or as strings.

```
str_sort(x = c("pear", "orange", NA, "berry", "nut"), decreasing = FALSE, na_last = TRUE, locale = "en", numeric = FALSE)

[1] "berry" "nut" "orange" "pear" NA
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

Use the functions demonstrated from the *stringr* package to complete the following.

1. Make a character vector with 26 elements "Letter: A", "Letter: B", "Letter: C" etc. Make use of `CharacterData004`.
2. For each of the following, make a character vector with a single element. The single element should be a string containing all of the elements of the indicated vector connected as indicated.
 - a. `CharacterData001`; link using spaces
 - b. `CharacterData002`; link using underscores
3. Split these strings each into a character vector with one word per element. Leave no extraneous spaces or characters.
 - a. `MoreCharacter003`
 - b. `MoreCharacter004`
 - c. `MoreCharacter005`
 - d. `MoreCharacter006`
4. Make a character vector with 10 elements. The elements should indicate Santa Claus' fictional enthusiasm, as he makes his fictional trip around the world. Each of the following lines should be an element.


```
Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!
Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!
Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!
Ho!Ho!Ho!Ho!Ho!Ho!Ho!Ho!
Ho!Ho!Ho!Ho!Ho!Ho!Ho!
Ho!Ho!Ho!Ho!Ho!
Ho!Ho!Ho!Ho!
Ho!Ho!Ho!
Ho!Ho!
Ho!
```
5. Make a character vector with 26 elements. The first element should be "I like the letter A.". The second element should be "I like the letter B." This should continue until the twenty sixth element "I like the letter Z."

Chapter 7: Logical Data

Logical Data takes in TWO values¹⁶:

1. TRUE (T)
2. FALSE (F)

Logical data can be used to encode the answer to a Yes/No question. Use TRUE for Yes, and FALSE for No, When using logical data, it must be entered in uppercase letters. Otherwise, R will not recognize it.

Collections of logical data can be held in an ordered (atomic) vector. A logical vector can be created using the **concatenate** function `c()`. Make sure to separate each logical data value with a comma.

```
c(TRUE, TRUE, FALSE, FALSE)
c(T, F, T, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F)
```

In fact, a single logical value on its own is a vector. It just happens to be a vector with one element.

Logical Operators

Computations using logical data can be performed.

The simplest operation is called **logical negation** or **NOT**. It turns a TRUE value into FALSE, and a FALSE value into TRUE.

```
!TRUE
[1] FALSE
!c(TRUE, FALSE)
[1] FALSE TRUE
```

Other logical operations take two arguments(inputs) and return a logical value. A logical AND determines if both inputs are TRUE. A logical OR determines if at least one input is TRUE. The **exclusive OR, xor()**, determines if exactly one of the inputs is TRUE.

```
# Logical AND &&
TRUE && TRUE
[1] TRUE

# Logical OR ||
TRUE || TRUE
[1] TRUE

xor(TRUE, T)
[1] FALSE
```

The operations indicated by **&&** and **||** will only look at one element vectors.

These operations can be used in sequence (taking the result of one operation and plugging it into another). You just need to enclose the first operation inside parentheses.

```
(TRUE && !TRUE) || TRUE
[1] TRUE
```

¹⁶ This is a slight lie. There is a third logical value that represents missing values

Vectorized Operations

If you want to perform these operations in a pairwise(element-wise) fashion, you want to use the **vectorized** version of these operators.

```
# Vectorized AND
c(TRUE, TRUE, FALSE, FALSE) & c(T, F, T, F)

[1] TRUE FALSE FALSE FALSE

# Vectorized OR
c(TRUE, TRUE, FALSE, FALSE) | c(T, F, T, F)

[1] TRUE TRUE TRUE FALSE

# exclusive Logical OR
xor(c(TRUE, TRUE, FALSE, FALSE), c(T, F, T, F))

[1] FALSE TRUE TRUE FALSE
```

As you can see, `xor()` will work with single element vectors or vectors with more elements.

Recycling

For operations with two inputs, if one input has fewer elements than the other, R will **recycle** (repeat) the shorter vector, until it is long enough to perform the vectorized operation.

```
c(TRUE, TRUE, FALSE, FALSE) & c(T, F)

[1] TRUE FALSE FALSE FALSE
```

In the first line, R repeated the vector with two elements. The second vector was treated like **c(T,F,T,F)**.

```
c(TRUE, TRUE, FALSE, FALSE) | c(T, F, T)

Warning in c(TRUE, TRUE, FALSE, FALSE) | c(T, F, T): longer object length is not a multiple of shorter object length

[1] TRUE TRUE TRUE TRUE
```

In this example, a warning is given. It indicates that the longer vector is not a multiple of the shorter vector, in terms of length. This is true. The longer vector had a length of 4, while the shorter had a length of three. In this case, the vector with length three was only partially recycled. To perform the computations, the second vector was treated like **c(T,F,T,T)**.

Any & All

Given a vector of logical data, you may wonder if **any** or **all** of the values are TRUE

```
any(c(F, T, F, F))

[1] TRUE

any(c(F, F, F, F))

[1] FALSE

all(c(T, T, T, T))

[1] TRUE

all(c(T, T, T, F))

[1] FALSE
```

One Interesting Fact

In certain contexts, R will treat logical data values as numbers. Namely, TRUE will be treated as one(1) and FALSE will be treated as zero(0).

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

Display any code you use to find your answer. In as few sentences as possible, state your answer.

1. Individually, determine the value of the following:
 - a. FALSE or TRUE
 - b. TRUE or TRUE
 - c. FALSE and TRUE
 - d. TRUE or FALSE
2. After creating two logical vectors, use a single line of code to determine the value of the following:
 - a. FALSE or TRUE
 - b. TRUE or TRUE
 - c. FALSE or TRUE
 - d. TRUE or FALSE
3. Translating these directly into the appropriate logical operations, determine which of these are TRUE:
 - a. (T or F) and T
 - b. (T and T) or (F and T)
 - c. not (not T or not F)
 - d. not not not TRUE
4. Compute the value of each without using recycling.
 - a. `c(T, T, F, F, TRUE)` and T
 - b. `not c(F, T, T)` and `c(F, T, F, T, F)`
 - c. `c(T, not FALSE, T, TRUE, T, T)` or `c(F, F, not T)`
 - d. `not c(FALSE, not FALSE, T, TRUE, T, T, FALSE)` or `not c(not F, not F, not FALSE, T)`
5. Using the results from question 4, determine if **any** and/or **all** of the resulting elements are TRUE.
6. Translating these directly into the appropriate logical operations, determine which of the following result in `c(T, T, T, T, T)`
 - a. `c(T, T, F, F, TRUE)` or F
 - b. `c(T, T, F, F, T, F)` or `not c(F, T, F)`
 - c. `c(T, not FALSE, T, TRUE, T, T)` and `not c(F, F, FALSE, not T)`
 - d. `not c(T, not FALSE, T, TRUE, T, T)` or `c(not F, not F, not FALSE, T)`

7. Perform a series of logical operations on the given vectors so that the result is a vector whose entries are all TRUE. All vectors in each part must be used at least once. Operations from the previous vectors can't be used.
(Experiment)

- a. $c(T, T, F, F)$, $c(T, F, T, F)$
- b. $c(T, T, T, T, F, F, F, F)$, $c(T, T, F, F)$, $c(T, F)$

Chapter 8: Comparisons

Given two data values (logical, numeric, character) it will be important to determine how they relate to each other. Are these values the same? Is one greater than the other? Does one come before the other?

The comparison operators can be used to answer these questions. They take any two data values, and return a logical value indicating the validity of the comparison.

As with many of the operators we have seen before, these operators are vectorized, and will recycle shorter vectors.

Ordering (<, <=, >=, >)

To order numeric values, we ask if one value is less than < or greater than > another value. If we want to include equality, add on an equal sign. Consider the possible comparisons between the numbers 0 and 1.

```
# Less than ?Comparison
0 < 1

[1] TRUE

1 < 0

[1] FALSE

# Less than or equal to
0 <= 1

[1] TRUE

1 <= 0

[1] FALSE

# greater than ( or equal to)
c(0, 1) > 1

[1] FALSE FALSE

c(0, 1, 1, 0) >= c(1, 0)

[1] FALSE TRUE TRUE TRUE
```

In the last line, realize that the $c(1,0)$ vector on the right hand side is being recycled. The comparisons that are being performed look (invisibly) like $c(0,1,1,0) \geq c(1,0,1,0)$.

These ordering operators are not limited to comparing numeric data only. They can be used on any of the data types. Strings will use an alphabetic ordering according to the locale setup in your computer.

Equality

Checking whether two values, or vectors, are the same can be tricky. The obvious choice for checking equality would be the equal sign =. However, the equal sign on its own performs a different function. The appropriate operator is a double equal sign ==.

```
7 == 14/2

[1] TRUE

c("a", "b") == c("a", "b", "a", "b")

[1] TRUE TRUE TRUE TRUE

c("a", "b") == c("a", "B", "a", "Horse")

[1] TRUE FALSE TRUE FALSE
```

The double equal sign will return a logical value, or vector of logical values. It is vectorized and will recycle a shorter vector if needed. As we can see, it will compare character vectors. It will also compare logical vectors. In each case, it compares corresponding elements.

The double equal sign does have some things to be aware of:

1. `==` looks for exact equality in the computers representation of a value.

Because of how a computer stores numbers, two values that mathematically are the same, may not be the same when stored. Consider the square root of 2.

```
sqrt(2) * sqrt(2) == 2L
```

```
[1] FALSE
```

```
sqrt(2) * sqrt(2)
```

```
[1] 2
```

Theoretically, this should be TRUE. A computer approximates the square root of 2. When this approximation is squared, a small amount error is included. It will approximate 2. As a remedy, we might ask if two values are almost equal or nearly equal. This can be done using the *all.equal()* function in base R or the *near()* function included in the *dplyr* package. Both of these functions have a small tolerance for some difference between two values.

```
all.equal(sqrt(2) * sqrt(2), 2)
```

```
[1] TRUE
```

```
dplyr::near(sqrt(2) * sqrt(2), 2)
```

```
[1] TRUE
```

Both functions can be used on vectors of values, but the outputs are very different. *all.equal()* will return TRUE if the difference in all pairs is within the given tolerance. However, it returns a string giving the average relative difference¹⁷.

```
all.equal(c(1.000000004, 2.000000009), c(1.000004, 2.00000008))
```

```
[1] "Mean relative difference: 1.332333e-06"
```

dplyr::near() will return a logical vector containing the results for each pairing.

```
dplyr::near(c(1.000000007, 2.000000009), c(1, 2.000000008), tol = 5e-09)
```

```
[1] FALSE TRUE
```

2. `==` does not compare the vectors as a whole. It compares corresponding elements

```
c("a", "b") == c("a", "b", "a", "b")
```

```
[1] TRUE TRUE TRUE TRUE
```

Allow for recycling of the shorter vector, the elements of these vectors match up. But, as a whole these vectors are not the same. Checking if any two R objects (not just vectors) are the same is done with the *identical()* function.

```
identical(c("a", "b"), c("a", "b", "a", "b"))
```

```
[1] FALSE
```

This function will determine if two elements are exactly the same. It will check many things not visible to the eye.

```
identical(2, 2L) #These are different types of numerical data.
```

```
[1] FALSE
```

```
identical(2, "2") #These are different types of data
```

```
[1] FALSE
```

```
2 == "2" #These are different types of data, but one is coerced into the other type.
```

```
[1] TRUE
```

¹⁷ What is average relative distance?

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Using the following sets of vectors, determine how elements from the first vector are related to the corresponding elements from the second. Use the indicated comparison.
 - a. 3, 3L (equality)
 - b. `c(1, 2, 3, 4, 5, 6)`, 4 (less than)
 - c. `c(1, 2, 3, 4, ..., 25)`, `c(sqrt(1), sqrt(2), ..., sqrt(25))^2`, (greater than or equal to)
 - d. `c(1, 2, 3, 4, ..., 25)`, `sqrt(c(1^2, 2^2, ..., 25^2))`, (less than or equal to)
2. Using the following sets of vectors, determine how elements from the first vector are related to the corresponding elements from the second. Use the indicated comparison.
 - a. "a", "A" (equality)
 - b. "a", "A" (less than)
 - c. `c("a", "b", "c")`, `c("b", "c", "d")` (greater than)
 - d. `c(1, 2, 3, 4)`, `c("1", "2", "3", "4")` (equality)
3. Using the following sets of vectors, determine how elements from the first vector are related to the corresponding elements from the second. Assign a variable to the result. Determine if **any** of the results are *TRUE*. Use the indicated comparison.
 - a. `c("the", "dude", "abides")`, `c("thematic", "duarchy", "aberatted")` (less than)
 - b. `c(1, "2", 4, 6)`, `c("one", "two", "three", "nine")` (greater than or equal)
 - c. `c(TRUE, FALSE)`, `c(1, 0)` (equality)
 - d. `c(TRUE, FALSE)`, `c("1", "0")` (equality)
4. Determine if the two given vectors are identical.
 - a. `c("the", "dude", "abides")`, `c("the", "dude", "abides", "the", "dude", "abides")`
 - b. `c("the", "dude", "abides", "the", "dude", "abides")`, `c("the", "dude", "abides", "the", "dude", "abides")`
 - c. `c(TRUE, FALSE)`, `c(1, 0)`
 - d. `c(1, 2, 3, 4, ..., 25)`, `sqrt(c(1^2, 2^2, ..., 25^2))`

5. Determine if the two given vectors are identical, and if **all** of the paired components are equal (==). If there is a discrepancy, give an explanation.
- a. 9, "9"
 - b. c("a", "aa", "aaa"), c("a", "aa", "aaa", "a", "aa")
 - c. c(1L, 2L, 3L, 4L, 5L, 6L, 7L), c(1L, 2, 3L, 4L, 5L, 6L, 7L)
 - d. c(T, T, F, T), c(TRUE, TRUE, FALSE, TRUE)

Chapter 9: Vectors

We have encountered four types of (atomic) vectors depending on the type of data contained in it.

1. Logical
2. Numerical - Integer/Floating Point
3. Character

Given a vector, the type of data contained in the vector can be determined using the `typeof()` function.

```
typeof(c(1, 2, 3, 4))
[1] "double"

typeof(c("A", "B", "String"))
[1] "character"
```

Vector Attributes

Besides the data explicitly contained in a vector, other pieces of information can be attached to the vector. Names can be given to each data value. Names can be added immediately as the data is being entered.

```
exampleData <- c(data1 = "apple", data2 = "orange", data3 = "pineapple", data4 = "strawberry")
exampleData

      data1      data2      data3      data4
"apple"    "orange" "pineapple" "strawberry"
```

If the names need to be attached, or updated, this can be done using the `names()` function. It will require a character vector with a name for each piece of data. The names should appear in the same order as the data they represent.

```
names(exampleData) <- c("newName1", "newName2", "newName3", "newName4")
exampleData

      newName1      newName2      newName3      newName4
"apple"        "orange"      "pineapple" "strawberry"
```

Comments can be added to describe the contents of the vector using the `comment()` function. Comments should be written as strings. The `comment()` function displays the comments also.

```
comment(exampleData) <- "This is a comment."
comment(exampleData)

[1] "This is a comment."
```

The `attributes()` function can be used to display a summary of these additional attributes.

```
attributes(exampleData)

$names
[1] "newName1" "newName2" "newName3" "newName4"

$comment
[1] "This is a comment."
```

Vector Length

The number of elements in a given vector can be determined with the `length()` function.

```
length(exampleData)

[1] 4
```

Numerical Sequences

A numerical vector containing a regular sequence of numbers is often needed. The most basic is a sequence of consecutive integers. This is generated with a colon separating the first and last integer in the sequence.

```
sequentialIntegers <- 3:8 #Non-Integers can be used, but the sequences increments by 1.
sequentialIntegers
```



```
[1] 3 4 5 6 7 8
```

For other regular sequences, you want to use the `seq()` function. It can be specified by three out of four arguments.

1. *from/to* - the starting / end values, at least one of these is needed.
2. *by* - the size of the increment between values
3. *length.out* - the number of values to be produced.

```
seq(from = 4, to = 6, length.out = 9)
seq(from = 4, to = 6.1, by = 0.25)
seq(from = 4, by = 0.25, length.out = 9)
seq(to = 6, by = 0.25, length.out = 9)
```

All of these produce the same result.

```
[1] 4.00 4.25 4.50 4.75 5.00 5.25 5.50 5.75 6.00
```

An integer sequence that from one to the length of an *R* object is useful. This can be done using colon notation, or `seq()`, along with the `length()` function. Alternatively, `seq_along()` accomplishes this with a single function.

```
seq_along(exampleData)
```

```
[1] 1 2 3 4
```

Replicating Vectors

Repeating patterns can be created with the `rep()` function. A vector containing the basic parts of the pattern is needed. The arguments control the method of replication.

```
givenVector <- c("a", 9)
rep(givenVector, times = 3)

[1] "a" "9" "a" "9" "a" "9"

rep(givenVector, each = 3, times = 2)

[1] "a" "a" "a" "9" "9" "9" "a" "a" "a" "9" "9" "9"

rep(givenVector, each = 3, times = 2, length.out = 10)

[1] "a" "a" "a" "9" "9" "9" "a" "a" "a" "9"
```

Combining Vectors

New vectors can be created by concatenation using variable names and explicitly written vectors.

```
x <- 7:4 #Sequential Vector
y <- rep(9, times = 3) #Replicated Vector
newVector <- c(x, y, 10, x)
newVector

[1] 7 6 5 4 9 9 9 10 7 6 5 4
```

This code chunk creates a vector *x* containing 7 down to 4, a vector *y* replicating the number 9 three times. Finally, the vector *newVector* is created with elements from *x*, then from *y*, followed by 10, and finishing with the elements from *x*.

Placeholder Vectors

When computations are repeated with different values, the results will need to be stored. If you know how many different sets of computations are to be performed, it can be useful to have a vector that can be filled in with the results.

```
resultsVector1 <- vector(mode = "logical", length = 5)
resultsVector2 <- rep(NA, length = 5) # NA stands for Not Available. It is used for missing values.
resultsVector1
resultsVector2

[1] FALSE FALSE FALSE FALSE
[1] NA NA NA NA NA
```

The mode in the `vector()` function can be replaced with integer, double, or character.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Determine the type of vector for each of the following: integer, double, character, logical
 - a. `c(1, 2L, 3L, 4L, 5L, 6L, 7L, 8L)`
 - b. `c(T, T, T, T, "F")`
 - c. `c(T, 1)`
 - d. `c("a", T, 10)`
2. For each set of attributes, create an appropriate vector. (Spaces are not allowed in names)
 - a. names, in order: one squared, two squared, three squared, four squared, ..., ten squared
 - b. comment: These are the fourth roots of the integers 2 through 10.
 - c. names, in order: lower a, upper B, lower c, upper D, ..., upper J comment: Letters a-j alternating case.
3. Create the following vectors and assign a variable to the result
 - a. `c(1, 2, 3, 4, ..., 1000)`
 - b. `c(sqrt(1), sqrt(2), sqrt(3), ..., sqrt(1000))`
 - c. `c(2, 4, 6, 8, ..., 1000)`
 - d. `c(1, 3, 5, 7, ..., 999)`
 - e. `c(10, 9, 8, 7, ... 0, -1, -2, ..., -50)`
4. Create the following vectors and assign a variable to the result
 - a. `c(1, 1.5, 2, 2.5, 3, 3.5, ..., 50)`
 - b. `c(2.2, 3.2, 4.2, 5.2, 6.2, ..., 100.2)`
 - c. `c(-1, -.6, -.2, .2, .6, 1.0, ..., 5.4)`
 - d. `c(ln(8.7), ln(9), ln(9.3), ln(9.6), ln(9.9), ln(10.2), ..., ln(15))`
5. Create a vector with the following properties:
 - a. Nine evenly spaced elements starting at -0.123 and ending at 68
 - b. Nineteen elements starting at 90 and decreasing by 12.
 - c. Elements starting at 9, increasing by .3, no larger than 21.2
 - d. Square root of all multiples of 10 between 1 and 2001.
 - e. Twenty-five more than the square root of all multiples of 82 between 1 and 2001.

6. Use the vectors listed below, or functions of them, to create the indicated vector:

VectorsA, VectorsGirl, VectorsOne, VectorsInts, VectorsZero

- a. `c(1, 4, 5, 6, 7)`
- b. A vector with 100 elements. The values alternate between "a" and "girl" with the first element being "a"
- c. `c("a", "1", "a", "girl", "4", "5", "6", "7", a)`
- d. A numerical vector with 200 elements. Each element is a 1.
- e. `c("girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl", "girl")`
- f. `c("girl", "girl", "girl", "girl", "girl", "a", "girl", "girl", "girl", "girl", "girl", "a", "girl", "girl", "girl", "girl", "girl", "a", "girl", "girl", "girl", "girl", "a")`
- g. `c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11)`
- h. `c(4, 5, 6, 7, 8, 10, 12, 14, 12, 15, 18, 21, 16, 20, 24, 28)`

7. Determine the number of elements in the following vectors.

- a. `seq(from = 156, to = 10000, by = pi)`
- b. `seq(from = 203, to = -56, by = -2.6)`
- c. `c(14.6, 15.1, 15.6, 16.1, ..., 200.1)`
- d. Elements start at -7.2, increasing by .3, no larger than 56.8

Chapter 10: Filtering & Subsetting Vectors

Given any type of vector with any number of elements, at times you will want to work with only smaller portion of the elements in the given vector. This can be done numerical values or with logical values.

Numerical Subsetting

To indicate which portion of a given vector you want to use, *R* needs to be given the **indexes** or **positions** of the desired elements need. The index is the numerical position as you count through the given vector.

```
givenVector <- c("a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k", "l", "m", "n", "o", "p", "q", "r", "s", "t", "a", "b",
"b", "c", "a")
givenVector

[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s" "t" "a" "b" "c" "a"
```

To determine the index of a value, count which element it is. *R* helps with this. At the head of each line, an integer is in square brackets. That integer indicates the index of the first element in that line. The capital T has an index of [20].

To select a single element from a vector, follow a straight forward syntax. Start with the name of the vector, followed by a set of square brackets that contain the index of the desired element.

```
# Selecting the 20th element
givenVector[20]

[1] "t"
```

A subset does not need to be restricted a single element. A numerical vector containing multiple values can be used.

```
# Selecting the 20th element
givenVector[c(1, 2, 20, 23, 24)]

selectThese <- c(1, 2, 20, 23, 24)
givenVector[selectThese]

[1] "a" "b" "t" "c" "a"
```

It may be that you want to select all elements from a given vector **except** for a given set. Use the same procedure as before, but use negative indexes

```
# Selecting the 20th element
givenVector[-c(1, 2, 20, 23, 24)]

givenVector[-selectThese]
```

Either of these will produce:

```
[1] "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s" "a" "b"
```

The *which()* function can be used to determine the indexes of all elements that satisfy a certain condition.

```
# Determine the index where 'a' elements are located.
which(givenVector == "a")

[1] 1 21 24

LOCATEa <- which(givenVector == "a")
givenVector[LOCATEa]

[1] "a" "a" "a"
```

Subsetting by Name

When names are assigned, a character vector inserted inside square brackets selects values.

```
exampleData <- c(data1 = 2.2, data2 = 9, data3 = -1)
exampleData[c("data2", "data3")]

data2 data3
9 -1
```

Reordering by filtering

The values in the index vector do not have to be arranged in an increasing order. They can be used to reorder all, or a subset of a vector.

```
givenVector <- c("z", "y", "x", "a", "b", "c")
givenVector[c(3, 2, 1)] #List the third, then second, then first value
[1] "x" "y" "z"
```

The `order()` function will reshuffle the indexes of all elements, so that the elements can be put in order

```
order(givenVector)
[1] 4 5 6 3 2 1

alphabeticalOrder <- order(givenVector)
givenVector[alphabeticalOrder]
[1] "a" "b" "c" "x" "y" "z"
```

Logical Subsetting

When using logic to subset/filter a vector, you must provide a *TRUE* or *FALSE* value for each element in the given vector.

```
givenVector <- 1:6
givenVector
[1] 1 2 3 4 5 6

selectThese <- c(TRUE, FALSE, TRUE, FALSE, TRUE, FALSE)
givenVector[selectThese]
[1] 1 3 5
```

Using this logical vector, only the odd numbers were selected from the numbers. The same result can be produced if we set up a comparison that will produce the same vector of *TRUE* and *FALSE* values. This comparison will determine if the remainder is 1 when each given value is divided by two.

```
givenVector%%2 == 1
[1] TRUE FALSE TRUE FALSE TRUE FALSE

selectOdd <- givenVector%%2 == 1
givenVector[selectOdd]
[1] 1 3 5
```

A more readable version of this can be seen if you use the `subset()` function.

```
# Argument \#1 is the vector to subset. Argument \#2 is the subset criteria.
subset(givenVector, givenVector%%2 == 1)
[1] 1 3 5
```

Replacing elements by filtering

When some elements in a vector need updating, filtering can be used to identify the elements, and then new values can be assigned.

```
dat <- c(1, 2, 3, NA, NA, 6, NA) # NA represents missing values in positions 4,5, and 7.
dat
[1] 1 2 3 NA NA 6 NA

dat[c(4, 5, 7)] <- c(104, 105, 107) # Replace the NA values
dat
[1] 1 2 3 104 105 6 107
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Extract the indicated element from the given vector.
 - a. The fifth element of **CharacterData004**.
 - b. The last element of **CharacterData005**.
 - c. The thirty-seventh element of **Filtering001**.
2. Extract the indicated elements from the given vector. Your answer for each should be a single vector with multiple elements. The elements should appear in the same order that as the vector they were extracted from. Each should be completed with a single line of code.
 - a. The tenth through the fifteen elements of **CharacterData001**.
 - b. The last five elements of **CharacterData006**.
 - c. The first twenty elements of **Filtering001**.
3. Extract the indicated elements from the given vector. The elements should appear in the same order that as the vector they were extracted from. Your answer for each should be a single vector with multiple elements. Each should be completed with at most two lines of code.
 - a. The elements in the odd indexes(positions) of **CharacterData004**
 - b. The elements of **CharacterData005** that have an index position that is a multiple of 4.
 - c. The elements **CharacterData001** whose index(position) is equal to a value in the Fibonacci Sequence of numbers.
4. Extract the indicated elements from the given vector. The elements should appear in the same order that as the vector they were extracted from. Your answer for each should be a single vector with multiple elements. Each should be completed with at most two lines of code.
 - a. The elements in **Filtering001** that are greater than or equal to 9.9.
 - b. The elements in **Filtering001** that are between 0 and 0.025.
 - c. The elements of **CharacterData001** that are either the string **if** or the string **this**.
5. For the indicated vector, determine the indices of the elements with the given properties.
 - a. The elements in **Filtering001** that are less than -9.95.
 - b. The elements of **CharacterData001** that start with a "W" or "w".
 - c. The elements of **CharacterData002** whose first character is a lower case a.
6. Extract the indicated elements from the given vector. The elements should appear in the same order that as the vector they were extracted from. Your answer for each should be a single vector with multiple elements. Each should be completed with at most two lines of code.
 - a. The elements of **CharacterData001** that are after the letter U in the dictionary.
 - b. The elements of **CharacterData002** whose second character is the letter M (either case).
 - c. The elements of **CharacterData002** whose start with a W followed by an H. (either case)
7. Create the following vectors. These should not be created by inspecting the indicated vectors.
 - a. A vector with two elements. The two elements should be the smallest and largest elements found in **Filtering001**. (Smallest, Largest)
 - b. A vector with four elements. It should contain the four middle values (in terms of size) of **Filtering001**. (Smallest --- Largest)

8. Create the following vectors. These should not be created by inspecting the indicated vectors.
- a. A vector with two elements. The two elements should be the smallest and largest elements found in **Filtering001**. (Smallest, Largest)
 - b. A vector with four elements. It should contain the four middle values (in terms of size) of **Filtering001**. (Smallest —– Largest)

Chapter 11: Matrices & Arrays

Matrices and Arrays are more general versions of vectors. If you think of a vector as a set of numbers written on a line, then think of a matrix as a set written out on a grid. If you start stacking matrices, then a three dimensional array is produced. However, from R point of view, a matrix or an array is a vector with a dimension attribute. The dimension information indicates how the values are indexed.

Matrices

To Construct a matrix, the **matrix()** function is used. Its first argument is the *data* to be included in the form of a vector. The next two arguments *nrow* and *ncol* indicate the number of rows and number of columns included. Only one if these is required, the other can be found by looking at the length of the data vector. The last argument *byrow* indicates how the matrix is filled in.

```
dat <- 1:24
mat.byrow.FALSE <- matrix(data = dat, nrow = 4, ncol = 6, byrow = FALSE)
mat.byrow.FALSE
```

| | [,1] | [,2] | [,3] | [,4] | [,5] | [,6] |
|------|------|------|------|------|------|------|
| [1,] | 1 | 5 | 9 | 13 | 17 | 21 |
| [2,] | 2 | 6 | 10 | 14 | 18 | 22 |
| [3,] | 3 | 7 | 11 | 15 | 19 | 23 |
| [4,] | 4 | 8 | 12 | 16 | 20 | 24 |

In this example, *byrow* is set to *FALSE*. Therefore, the columns are filled first from top to bottom, then moving left.

Looking at the margins of the matrix, the square bracketed values indicate the two part index for each value. An individual value is indicated by its row index, followed by a column index. The values should be separated by a comma. Vectors of multiple values can be used to filter a matrix.

```
# Values found in rows 1 and 4 AND columns 2, 3, 4, and 5.
mat.byrow.FALSE[c(1, 4), 2:5]
```

| | [,1] | [,2] | [,3] | [,4] |
|------|------|------|------|------|
| [1,] | 5 | 9 | 13 | 17 |
| [2,] | 8 | 12 | 16 | 20 |

```
# No values from rows 2 or 3 are include. Neither are values from column 1.
```

If all values in a subset of rows are desired, the column indices don't need to be specified. However, the comma can not be omitted. If a subset of columns is desired, the row values can be omitted.

```
# Rows 1 and 4, all columns
mat.byrow.FALSE[c(1, 4), ]
```

| | [,1] | [,2] | [,3] | [,4] | [,5] | [,6] |
|------|------|------|------|------|------|------|
| [1,] | 1 | 5 | 9 | 13 | 17 | 21 |
| [2,] | 4 | 8 | 12 | 16 | 20 | 24 |

```
# Columns 2 and 3, all rows
mat.byrow.FALSE[, c(2, 3)]
```

| | [,1] | [,2] |
|------|------|------|
| [1,] | 5 | 9 |
| [2,] | 6 | 10 |
| [3,] | 7 | 11 |
| [4,] | 8 | 12 |

If the *byrow* argument is set to *TRUE*, the rows are filled from left to right, starting with the top row and moving down. A final argument *dimnames* allows names to be attached to the rows and/or columns. The names need to be added as a list. Lists have not been described yet, they will be in the future.

```
rowlabel <- c("row1", "row2", "row3", "row4")
collabel <- c("column1", "column2", "column3", "column4", "column5", "column6")
mat.labels <- list(row = rowlabel, col = collabel)
dat <- 1:24
mat.byrow.TRUE <- matrix(data = dat, nrow = 4, ncol = 6, byrow = TRUE, dimnames = mat.labels)
mat.byrow.TRUE
```

| | col | | | | | |
|------|---------|---------|---------|---------|---------|---------|
| row | column1 | column2 | column3 | column4 | column5 | column6 |
| row1 | 1 | 2 | 3 | 4 | 5 | 6 |
| row2 | 7 | 8 | 9 | 10 | 11 | 12 |
| row3 | 13 | 14 | 15 | 16 | 17 | 18 |
| row4 | 19 | 20 | 21 | 22 | 23 | 24 |

To use the row or column names to filter elements from a matrix, use two character vectors whose elements list the names of the desired rows and columns.

```
mat.byrow.TRUE[c("row1", "row4"), c("column2", "column3", "column4", "column5")]
```

| | col | | | |
|------|---------|---------|---------|---------|
| row | column2 | column3 | column4 | column5 |
| row1 | 2 | 3 | 4 | 5 |
| row4 | 20 | 21 | 22 | 23 |

As it was said earlier, a matrix is a vector with an additional dimension attribute. The dimension of a matrix is a vector with the number of rows followed by the number of columns. Both matrices that were constructed have 4 rows and 6 columns.

```
dim(mat.byrow.FALSE)
```

```
[1] 4 6
```

```
dim(mat.byrow.TRUE)
```

```
[1] 4 6
```

Looking at the structure of either of these, you will see the vector of data listed as a vector with the indices running from 1 to 4 for the rows and 1 to 6 for the columns as well as an explicit indication of the dimension.

```
str(mat.byrow.TRUE)
```

```
int [1:4, 1:6] 1 7 13 19 2 8 14 20 3 9 ...
- attr(*, "dimnames")=List of 2
..$ row: chr [1:4] "row1" "row2" "row3" "row4"
..$ col: chr [1:6] "column1" "column2" "column3" "column4" ...
```

As with a vector, a matrix can only hold one type of data: logical, numeric, or character. From the method of construction, this should be clear. A matrix is built from a vector, which can only hold one type of data.

Arrays

A matrix is a two dimensional example of an array. Arrays can have as many dimensions as you like. However, after three dimensions, its impossible to picture their structure in higher dimensions. A three dimensional array can be pictured as at least two $rx c$ matrices stacked on top of each other. With this in mind, each element would be described by its row and column in one of these matrices. A third value would be needed to describe which matrix in the stack.

Creating an array, is similar to creating a matrix. A vector of data values is needed. This is passed to the `array()` function. A secondary argument `dim` indicates the three dimensions. Each matrix is filled from the first row to the last, starting with the first column and moving left. Once the first matrix is filled, the second matrix in the stack if filled in the same way. Care must be taken to make sure the original data is in the desired order.

```
dat <- c(1:20, 101:120, 1001:1020)
arr <- array(data = dat, dim = c(4, 5, 3))
arr

, , 1

      [,1] [,2] [,3] [,4] [,5]
[1,]    1    5    9   13   17
[2,]    2    6   10   14   18
[3,]    3    7   11   15   19
[4,]    4    8   12   16   20

, , 2

      [,1] [,2] [,3] [,4] [,5]
[1,]  101  105  109  113  117
[2,]  102  106  110  114  118
[3,]  103  107  111  115  119
[4,]  104  108  112  116  120

, , 3

      [,1] [,2] [,3] [,4] [,5]
[1,] 1001 1005 1009 1013 1017
[2,] 1002 1006 1010 1014 1018
[3,] 1003 1007 1011 1015 1019
[4,] 1004 1008 1012 1016 1020
```

When displayed, each 'stack' is indicated by ', , stack number'. With enough data, this can become very unmanagable to look through.

Filtering elements from an array works as you would for a matrix. In three dimensions, you list row, column, and then the stack. Leaving out a set of indices, indicates that all indices are desired. Depending upon the selection, *R* may return a vector, a matrix, or an array. Whichever form of data is simplest.

```
arr[1:2, 2:4, c(1, 3)] # An Array

, , 1

      [,1] [,2] [,3]
[1,]    5    9   13
[2,]    6   10   14

, , 2

      [,1] [,2] [,3]
[1,]  1005 1009 1013
[2,]  1006 1010 1014

arr[, , 3] # A matrix

      [,1] [,2] [,3] [,4] [,5]
[1,]  1001 1005 1009 1013 1017
[2,]  1002 1006 1010 1014 1018
[3,]  1003 1007 1011 1015 1019
[4,]  1004 1008 1012 1016 1020

arr[1:3, 4, 2] # A vector

[1] 113 114 115
```

The `str()` and `dim()` functions play the same roll they did for matrices. Names can be added, if desired. However, they could become very difficult to work with.

```
str(arr)

int [1:4, 1:5, 1:3] 1 2 3 4 5 6 7 8 9 10 ...

dim(arr)

[1] 4 5 3
```

As with a vector and a matrix, an array can only hold one type of data: logical, numeric, or character.

An example of a four dimensional array is given below. Think of it as at least two three dimensional arrays stacked.

```
dat1 <- c(1:20, 101:120, 1001:1020)
dat2 <- c(21:40, 221:240, 2021:2040)
arr.4d <- array(data = c(dat1, dat2), dim = c(4, 5, 3, 2))
arr.4d

, , 1, 1

      [,1] [,2] [,3] [,4] [,5]
[1,]     1     5     9    13    17
[2,]     2     6    10    14    18
[3,]     3     7    11    15    19
[4,]     4     8    12    16    20

, , 2, 1

      [,1] [,2] [,3] [,4] [,5]
[1,]   101   105   109   113   117
[2,]   102   106   110   114   118
[3,]   103   107   111   115   119
[4,]   104   108   112   116   120

, , 3, 1

      [,1] [,2] [,3] [,4] [,5]
[1,]  1001  1005  1009  1013  1017
[2,]  1002  1006  1010  1014  1018
[3,]  1003  1007  1011  1015  1019
[4,]  1004  1008  1012  1016  1020

, , 1, 2

      [,1] [,2] [,3] [,4] [,5]
[1,]    21    25    29    33    37
[2,]    22    26    30    34    38
[3,]    23    27    31    35    39
[4,]    24    28    32    36    40

, , 2, 2

      [,1] [,2] [,3] [,4] [,5]
[1,]   221   225   229   233   237
[2,]   222   226   230   234   238
[3,]   223   227   231   235   239
[4,]   224   228   232   236   240

, , 3, 2

      [,1] [,2] [,3] [,4] [,5]
[1,]  2021  2025  2029  2033  2037
[2,]  2022  2026  2030  2034  2038
[3,]  2023  2027  2031  2035  2039
[4,]  2024  2028  2032  2036  2040
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Create the following **matrix** that is filled with **letters**:

```
[,1] [,2] [,3] [,4] [,5] [,6]
[1,] "a"  "b"  "c"  "d"  "e"  "f"
[2,] "g"  "h"  "i"  "j"  "k"  "l"
[3,] "m"  "n"  "o"  "p"  "q"  "r"
[4,] "s"  "t"  "u"  "v"  "w"  "x"
```

2. Create the matrix displayed below.

```
[,1] [,2] [,3] [,4] [,5] [,6] [,7]
[1,] 1   9   3   12  5   15  7
[2,] 3   12  5   15  7   1   9
[3,] 5   15  7   1   9   3   12
[4,] 7   1   9   3   12  5   15
```

3. Extract the fourth row of the matrix **fruitMatrix**.
4. Extract the fifth column of the matrix **fruitMatrix**.
5. Create a new matrix by extracting a portion of the matrix **fruitMatrix**. It should consist of the elements of **fruitMatrix** that are in middle four rows of **fruitMatrix** AND the middle three columns of **fruitMatrix**.
6. Create a new vector by extracting a portion of the matrix **fruitMatrix**. It should consist of the elements in the fifth column of **fruitMatrix** that are also in the second, third, and fifth row of **fruitMatrix**.
7. Create a new matrix by extracting some columns from the matrix **numberMatrix**. The columns that are extracted should be those whose last element are between 15(included) and 60(excluded).
8. Create a new matrix by extracting some rows from the matrix **numberMatrix**. The rows that are extracted should be those whose first element is less than or equal to 25.

This page left intentionally blank for more problems

Chapter 12: Constants Values and Vectors

Several useful values and vectors are built into the base *R* package. These relieve you from needing to type them in manually, or loading them from outside source.

Pi

The numerical constant π can be obtained by using the *pi* variable.

```
pi
[1] 3.141593
```

Letters

The 26 letters of the Roman alphabet are loaded in either upper or lower case form into two vectors.

```
letters
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s" "t" "u" "v" "w" "x" "y" "z"

LETTERS
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q" "R" "S" "T" "U" "V" "W" "X" "Y" "Z"
```

Months

The months of the year are also accesible.

```
month.abb
[1] "Jan" "Feb" "Mar" "Apr" "May" "Jun" "Jul" "Aug" "Sep" "Oct" "Nov" "Dec"

month.name
[1] "January" "February" "March" "April" "May" "June" "July" "August" "September" "October" "November" "December"
```

Sometimes, it is useful to have a larger collections of strings to experiment on. By loading the *stringr* package you are given access to several interesting

Fruit, Words, and Sentences

fruit gives access to a character vector containing 80 names of fruit.

```
library(stringr)
fruit[1:10]
[1] "apple" "apricot" "avocado" "banana" "bell pepper" "bilberry" "blackberry" "blackcurrant" "blood orange" "blueberry"
```

words gives access to a character vector containing 980 words.

```
words[1:10]
[1] "a" "able" "about" "absolute" "accept" "account" "achieve" "across" "act" "active"
```

Sentences

`sentences` gives access to a character vector containing 980 words.

```
sentences[1:10]
```

```
[1] "The birch canoe slid on the smooth planks." "Glue the sheet to the dark blue background." "It's easy to tell the depth of a well." "These days a chicken leg is a rare dish." "Rice is often served in round bowls."
[6] "The juice of lemons makes fine punch." "The box was thrown beside the parked truck." "The hogs were fed chopped corn and garbage." "Four hours of steady work faced us." "A large size in stockings is hard to sell."
```

A Lorem Ipsum Generator

The `stri_rand_lipsum()` is a Lorem Ipsum generator included in the `stringi` package. It will produce a string of dummy text. Its main purpose is to produce text that is useful for considering page layouts, without being distracted by the actual content.

```
library(stringi)
stri_rand_lipsum(n_paragraphs = 3, start_lipsum = TRUE)
```

```
[1] "Lorem ipsum dolor sit amet, odio ipsum arcu posuere nec sit tristique. Proin at ac maecenas platea libero, sapien at mi. Sed orci a enim. Tristique, ac, mus per magna mauris nec nascetur et porta nec faucibus. Id at, nascetur sodales id vitae vel dapibus. Metus mi, magna nunc eu nulla blandit. Faucibus tristique sapien curae sed hendrerit eu. Maximus, integer ac per accumsan odio ex id nulla ut vel facilisis. Eros, quisque netus congue aenean ut senectus."
[2] "Mauris blandit in sed ipsum, at leo tincidunt. Primis, luctus facilisi nisl, facilisi taciti curae. Gravida per, amet purus sem at tempus sociosqu non litora felis inceptos aptent facilisis. Ac donec tellus turpis congue. Risus, a eu in magnis ultricies nulla nibh. Metus vitae diam amet, curae pulvinar dolor ipsum sit. Erat aliquam eget urna congue magna ut, purus phasellus dictum euismod. Vitae ante cras vulputate tellus eu. Velit sodales nullam mauris nunc tortor curae in feugiat sit sit elementum, sed."
[3] "Felis rutrum, odio blandit aenean ultricies, diam! Proin euismod mus sed molestie ac at, pellentesque duis nam nisi. Orci lectus hendrerit ligula eros accumsan. Tempor sociosqu per fringilla nisl duis vitae vestibulum fusce vel etiam ut porttitor, sed ad cubilia. Scelerisque donec sed amet placerat, aptent libero non sed. Eros in pretium nec velit, nisi. Commodo quam commodo dictum aliquam fusce at curabitur habitant. Sed ut nunc sed in massa, rhoncus integer blandit. Dapibus id a elit in. In bibendum quis curabitur sit lacus ante. Lorem purus condimentum phasellus arcu vehicula mi lobortis commodo dui dui."
```

The first argument determines the number of paragraphs to generate. The second indicates whether the text should start with 'Lorem ipsum dolor sit amet'.

Colors

While not a vector, the `colors()` function returns vector containing all the named colors in *R*. When producing graphics, these names can be used to set the color of a graphical object. There are 657 available to choose from.

```
clrs <- colors()
clrs[1:20]
```

| | | | | | | | |
|---------------|---------------|----------------|-----------------|-----------------|-----------------|-----------------|--------------|
| [1] "white" | "aliceblue" | "antiquewhite" | "antiquewhite1" | "antiquewhite2" | "antiquewhite3" | "antiquewhite4" | "aquamarine" |
| "aquamarine1" | "aquamarine2" | "aquamarine3" | "aquamarine4" | "azure" | "azure1" | "azure2" | |
| [16] "azure3" | "azure4" | "beige" | "bisque" | "bisque1" | | | |

Time Zones

Like the `colors()` function, the `OlsonNames()` function returns a vector. This vector returns a vector of time zones available in *R*.

```
tz <- OlsonNames()
tz[1:10]
```

| | | | | | |
|----------------------|-----------------|----------------------|------------------|-----------------|-----------------|
| [1] "Africa/Abidjan" | "Africa/Accra" | "Africa/Addis_Ababa" | "Africa/Algiers" | "Africa/Asmara" | "Africa/Asmera" |
| "Africa/Bamako" | "Africa/Bangui" | "Africa/Banjul" | "Africa/Bissau" | | |

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

These should not be completed by inspecting the given vectors and typing in the answer.

1. What is the cube root of the fifth power of pi?
2. Make a character vector with a single element. The element should be a string with all the letters of the English alphabet in reverse alphabetical order. Both upper case and lower case letters should appear. Each lower case letter should be immediately followed by the upper case version. Letters should not be individually typed.
3. For each fruit that begins with the letter B in the **fruit** vector, determine how many characters are in its name. The result should be a single numerical vector.
4. TRUE/FALSE There is at least one sentence in **sentences** whose last letter is the letter z(either case). The result should be a logical vector of length one.
5. TRUE/FALSE All the words in **words** are at most 11 characters long. The result should be a logical vector of length one.
6. How many Olson time zones indicate in their text Europe?
7. For each fruit that ends with the letter r in the **fruit** vector, determine how many characters are in its name. The result should be a single numerical vector.
8. TRUE/FALSE There is at least one sentence in **sentences** whose last two letters, in any order are r and e (either case). The result should be a logical vector of length one.
9. TRUE/FALSE All the words in **words** are at least 2 characters long. The result should be a logical vector of length one.
10. How many Olson time zones indicate in their text **Greenwich Mean Time**?

This page left blank for more questions to be added in the future.

Chapter 13: Infinity and Undefined Values

Some computations result in outcomes that are too large to *R* to store or are undefined mathematically.

Infinity

In situations, the result of a computation is larger (in absolute value) than can be stored. *R* will store this value as $\pm Inf$, either a positive or negative infinity, depending on the case.

```
2^9999
[1] Inf
-10^1000
[1] -Inf
```

While dividing a non-zero value by zero is undefined in some contexts. It is often useful to assign an infinite value to this result. Depending on the numerator, *R* will indicate this with either a positive or negative infinity.

```
1/0
[1] Inf
-pi/0
[1] -Inf
```

As far as *R* is concerned, these infinities are numerical values. Standard arithmetic operations can be performed with them, and a consistent result will be given.

```
typeof(Inf)
[1] "double"
Inf + 1
[1] Inf
10/Inf
[1] 0
```

Testing for an infinity can be done a few ways. The `==` operator will detect either a positive or negative infinity. `is.infinite()` will detect either type.

```
x <- c(1, -1)/0
x == Inf
[1] TRUE FALSE
x == -Inf
[1] FALSE TRUE
is.infinite(x)
[1] TRUE TRUE
```

Not a Number

Some operations that are undefined should not produce a meaningful numeric result. This can occur when you are subtracting infinite values, or multiplying/dividing infinite values and zero. When the result should not be a meaningful numeric result, *R* indicates that the result is 'Not a Number'.

```
Inf - Inf
```

```
[1] NaN
```

```
Inf * 0
```

```
[1] NaN
```

```
0/0
```

```
[1] NaN
```

```
Inf/Inf
```

```
[1] NaN
```

While *NaN* is not a meaningful numerical value, *R* will still allow arithmetic operations to be performed on it. The result will also be *NaN*.

```
typeof(NaN)
```

```
[1] "double"
```

```
NaN + 1
```

```
[1] NaN
```

However, identifying a *NaN* value can not be done using the usual `==` double equal sign. Instead, the `is.nan()` function must be used.

```
NaN == NaN
```

```
[1] NA
```

```
is.nan(NaN)
```

```
[1] TRUE
```

If you want to identify *NaN* values or missing values *NA*, `is.na()` can be used.

```
NaNNA <- c(NaN, NA, 2)
```

```
is.na(NaNNA)
```

```
[1] TRUE TRUE FALSE
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

These should not be completed by inspecting the given vectors and typing in the answer.

1. How many **PossiblyInfinite** values are infinite?
2. How many **PossiblyNAN** values are not a number?

This page left blank for more questions to be added in the future.

Chapter 14: Missing Values & Vectors

When a data value is missing for some reason, it still needs to be recorded that some values should be listed. In *R*, *NA* can be used to indicate a missing data value. It stands for ‘Not Available’.

```
NA
[1] NA
typeof(NA)
[1] "logical"
```

As you can see, *NA* on its own is actually a third type of logical value. The other two being *TRUE* and *FALSE*. However, when *NA* is used in vector containing some other data type, it will be coerced/changed into the appropriate data type, be it numeric or character data.

```
givenVector <- c(1, 2, 3, NA, NA)
typeof(givenVector[4])
[1] "double"
```

Operations with Missing Values

Performing operations with *NA* values can be tricky. The *NA* value will infect the operations, and generally lead to a result of *NA*. This is something to be aware of.

```
givenVector + 1
[1] 2 3 4 NA NA
0 * NA
[1] NA
```

Many functions will have an logical argument *na.rm*. If this argument is set to *TRUE*, missing values will be ignored.

Detecting Missing Values

The `==` double equal sign operator can’t be used to detect *NA* values. It will return a value of *NA*.

```
NA == NA
[1] NA
```

As a workaround, two functions are provided that can logically check for *NA* values. *is.na()* is a vectorized function that will check each element of a vector to see if it is *NA* and return *TRUE/FALSE* for each. *anyNA()* will return a single logical value. If at least one value is *NA*, *TRUE* will be returned.

```
is.na(NA)
[1] TRUE
is.na(givenVector)
[1] FALSE FALSE FALSE TRUE TRUE
anyNA(givenVector)
[1] TRUE
```

Replacing Missing Values

Certain situation require replacing *NA* values with another value. If all *NA* values will be replaced with the same value, the replacement is done with a short line of code.

```
# Replace the NA values in givenVector with -100.
givenVector

[1] 1 2 3 NA NA

givenVector[is.na(givenVector)] <- -100
givenVector

[1] 1 2 3 -100 -100
```

The first line of code displays `givenVector`. The last two elements are *NA*. The second line of code does the replacement. As it is written, it can be tricky to read. `is.na(givenVector)` produces a vector of *TRUE/FALSE* values. Placing that vector inside the square brackets filters the values in `givenVector` corresponding to *TRUE*. Finally, the assignment operator assigns -100 to those two values.

NULL

NA is used to indicate that a value within a vector is absent. The *NULL* value is used to indicate the absence of a vector. Sometimes, *NULL* will be the result of an undefined computation. Sometimes, *NULL* will be used to erase some information.

```
length(NULL)

[1] 0

typeof(NULL)

[1] "NULL"
```

Detecting a *NULL* value is done with the `is.null()` function.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Determine if the given vectors contain any missing values. Your answer should be a single element logical vector for each. This should be indicated by a TRUE value if the given vector contains any missing values. Otherwise, FALSE. This should not be done by inspection.
 - a. MissingValues01
 - b. MissingValues02
 - c. MissingValues03
2. Determine if the given vectors contain any missing values. Your answer should be a single element logical vector for each. This should be indicated by a TRUE value if the given vector contains any missing values. Otherwise, FALSE. This should not be done by inspection.
 - a. firstNames
 - b. lastNames
 - c. idNumbers
3. Determine the indices of all missing values in the given vectors. Your answer should be a single vector for each of the following.
 - a. MissingValues01
 - b. MissingValues02
 - c. MissingValues03
 - d. firstNames
 - e. lastNames
 - f. idNumbers
4. Create a vector that contains all the non-missing values of the vector indicated below. (In order to not cause problems in future questions, the new vector should have the same name as the vector indicated below with the letter X appended to the end.)
 - a. MissingValues01
 - b. MissingValues03
 - c. firstNames
 - d. lastNames
5. For each pair of vectors, display the elements of the first vector that correspond to the missing elements of the second vector.
 - a. firstNames; lastNames
 - b. lastNames; firstNames
 - c. idNumbers; firstNames
 - d. idNumbers; lastNames
 - e. lastNames; idNumbers
 - f. A vector of both corresponding first and last names, last name first, seperated by a comma; idNumbers
6. For each vector below, replace the missing values with the indicated quantity/string. Do not display the result.
 - a. firstNames; Kobalt
 - b. lastNames; 000000

c. idNumbers; G-Ma

Waiting for more problems

Chapter 15: Coercion

The data type of a vector is determined by the values entered into it. However, **R** has two methods for coercing (changing) the data type of a vector: implicit and explicit coercion.

Implicit coercion occurs when a vector is created using a mixture of data types. In situations where an attempt is made to mix data types into a single vector, **R** will coerce the different data types into a common data type. A hierarchy exists for this implicit coercion with logical data at the bottom, character data at the top, and numeric data in the middle. When data types are mixed, lower types will move as high up as possible.

Explicit coercion occurs when a function is applied to a vector to change its data type. This does not need to be a vector that was created with mixed data types. Explicit coercion can be used to move vectors either up or down the hierarchy. However, care must be taken as vectors are moved down the hierarchy as some values can not be coerced 'down' and will be converted into *NA* values.

Implicit Coercion - Mixed Data Types

Numeric Data

As numeric data falls in the middle of the hierarchy, when mixed with other types it will either remain its own type or become character data. In this example, several numeric values are written into a vector that contains a character element.

```
Numeric_Char <- c("STRING", 0, 1, 2)
Numeric_Char

[1] "STRING" "0"      "1"      "2"
```

Each numerical value in the output is now surrounded by quotation marks. This is an indication that each numeric value is regarded as a single character now. These characters have lost their numeric value and mathematical operations can no longer be performed on them. Errors will result.

```
"0" * 1

Error in "0" * 1: non-numeric argument to binary operator
```

Logical Data

At the bottom of the hierarchy, logical data can be coerced into either numeric or character data depending upon what other data types are mixed in a given vector. This mix also affects the output.

A mix of logical and numeric data will become numeric data. The coercion process will turn TRUE values into the number one and FALSE values into the number zero.

```
Logical_Num <- c(9, TRUE, T, FALSE, F)
Logical_Num

[1] 9 1 1 0 0

typeof(Logical_Num)

[1] "double"
```

Additionally, if mathematical operations are performed on a set of vectors that mix such values, the same numerical replacement will occur.

```
T + 10

[1] 11
```

Mixing Logical data with Character data will result in a direct conversion into character data, even if numerical values are present.

```
Logical_Char <- c("STRING", TRUE, T, FALSE, F, 10)
Logical_Char

[1] "STRING" "TRUE"   "TRUE"   "FALSE"  "FALSE"  "10"
```

The existence of a character string determines that the resulting vector will be a character vector. The logical values are coerced into string values of either “TRUE” or “FALSE” even if *T* and *F* are used¹⁸.

Explicit Coercion

With the use of additional functions, the specifics of the coercion can be better controlled.

Becoming Character Data

The *as.character()* function can be used to explicitly coerce a vector into a character data.

```
allNumeric <- as.character(c(1, 2, 3))
allLogical <- as.character(c(TRUE, FALSE, T, F))
allNumeric

[1] "1" "2" "3"

allLogical

[1] "TRUE" "FALSE" "TRUE" "FALSE"
```

With the surrounding quotes, it should be clear that these vectors contain character data even though they were written as numeric/logical withing the *as.character()* function.

Becoming Numeric Data

The *as.numeric()* function can be used to explicitly coerce a vector into a numeric vector. *as.numeric()* will coerce numeric characters into numeric values. This allows mathematical computations to be performed on them. However, strings that can not be interpreted as numeric values will be converted in to *NA* values¹⁹.

```
Char_to_Num <- as.numeric(c("4", "-100.2", "Ten"))
Char_to_Num

[1] 4.0 -100.2 NA
```

Similarly, *as.numeric* can be used to explicitly coerce logical data into numeric data. The numeric conversion of TRUE becoming a one and FALSE becoming a zero occurs.

```
Logical_to_Num <- as.numeric(c(T, F, TRUE, FALSE))
Logical_to_Num

[1] 1 0 1 0
```

Becoming Logical Data

The *as.logical()* function can be used to explicitly coerce a vector into a logical vector data. As with *as.numeric()*, only certain values can be coerced into logical values. All others will become *NA* values.

```
Char_to_Logi <- as.logical(c("TRUE", "FALSE", "T", "F ", "0", "1", "Character"))
Char_to_Logi

[1] TRUE FALSE TRUE NA NA NA NA
```

The only coercible strings are “TRUE”, “FALSE”, “T”, “F”²⁰. Adding even so much as a space to any of these will result in an *NA* value as is clear from the string “F”.

When using *as.logical()* on a numeric vector, every numeric value will be converted into TRUE except for the value zero. Zeros will be converted into FALSE.

```
Num_to_Logi <- as.logical(c(0, 1, 10.9, -2))
Num_to_Logi

[1] FALSE TRUE TRUE TRUE
```

¹⁸ Notice that the presense of the numeric value 10 does not cause the logical values to become zeros and ones before becoming character strings.

¹⁹ When this happens, a warning can be displayed

²⁰ “True”, “true”, “False”, and “false” will also be coerced to TRUE and FALSE.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Coerce the indicated vector into the indicated type. Additionally, produce a logical value where it is TRUE if the resulting vector contains missing values (FALSE otherwise).
 - a. Coerce01; logical
 - b. Coerce02; logical
 - c. Coerce03; character
 - d. Coerce04; logical
 - e. Coerce01; numerical
 - f. Coerce05; numerical
2. Coerce the indicated vector into the indicated type in such a way that missing values are not introduced.
 - a. Coerce01; logical
 - b. Coerce02; logical
 - c. Coerce03; character
 - d. Coerce04; logical
 - e. Coerce04; numerical
 - f. Coerce05; character with numeric characters.
3. Using some/all of vectors Coerce01, Coerce02, and Coerce03 (at most once) create a vector with at least six elements of the indicated type. (Implicit coercion only should be used.) The vectors should be used with their names in increasing order.
 - a. character
 - b. numeric
 - c. logical
4. Using some/all of vectors Coerce01, Coerce02, and Coerce03 (at most once) create a vector with at most three elements of the indicated type. Explicit coercion should be used. The resulting vector should not have any missing values.
 - a. character
 - b. numeric
 - c. logical
5. Combine three different vectors whose name starts **Coerce** in such a way that the resulting vector does not have any missing values and is of the following type. Both implicit and explicit coercion can be used. Display the vector you created and its type. The vectors should be used with their names in increasing order.
 - a. character
 - b. numeric
 - c. logical

More to come

Chapter 16: Lists

Lists provide another way to store data. Lists differ from vectors in that lists can hold many different types of data.

Making a List

As a basic example, three vectors are created, each of a different type and different length.

```
numbers <- 1:10
characters <- letters[1:6]
logicals <- rep(c(T, F), times = 2)
```

To create a list containing these three vectors, use the `list()` function and type in the names of the created vectors. As *R* creates the list, it actually coerces/changes these vectors into sublists.

```
newList <- list(numbers, characters, logicals)
newList

[[1]]
[1] 1 2 3 4 5 6 7 8 9 10

[[2]]
[1] "a" "b" "c" "d" "e" "f"

[[3]]
[1] TRUE FALSE TRUE FALSE
```

As you examine printed list, you will see two sets of indexes. The double square brackets `[[]]` indicate which sublist being displayed. The single bracket `[]` are used to index the vectors themselves.

List Attributes

Besides the data explicitly contained in a list, other pieces of information can be attached to it. Names can be given to each sublist. This can be done in two different ways.

Names can be added, modified, and retrieved, and using the `names()` function²¹. It requires a character vector with a name for each sublist. The names appearing in the same order as the sublists.

```
names(newList) <- c("LN", "LC", "LL")
```

Alternatively, the names can be added as the list is being created. Notice that the `[[]]` are replaced by dollar sign and a name with a named list.

```
namedList <- list(LN = numbers, LC = characters, LL = logicals)
namedList

$LN
[1] 1 2 3 4 5 6 7 8 9 10

$LC
[1] "a" "b" "c" "d" "e" "f"

$LL
[1] TRUE FALSE TRUE FALSE
```

Comments can be added to describe the contents of the list using the `comment()` function. Comments should be written as strings. The `comment()` function displays the comments also.

```
comment(namedList) <- "This is a comment on a list."
comment(namedList)

[1] "This is a comment on a list."
```

The `attributes()` function can be used to display a summary of these additional attributes.

```
attributes(namedList)

$names
[1] "LN" "LC" "LL"
```

²¹ The same as is can with vectors.

```
$comment
[1] "This is a comment on a list."
```

List Structure

When trying to access the information in a list, it is useful to look at its underlying structure using the `str()` function. The `str()` function will list the names of each component as well as a simplified summary of the information contained in that component.

```
str(newList)

List of 3
 $ LN: int [1:10] 1 2 3 4 5 6 7 8 9 10
 $ LC: chr [1:6] "a" "b" "c" "d" ...
 $ LL: logi [1:4] TRUE FALSE TRUE FALSE

str(namedList)

List of 3
 $ LN: int [1:10] 1 2 3 4 5 6 7 8 9 10
 $ LC: chr [1:6] "a" "b" "c" "d" ...
 $ LL: logi [1:4] TRUE FALSE TRUE FALSE
- attr(*, "comment")= chr "This is a comment on a list."
```

The `str()` function is very useful with extracting information from lists. It can be very easy to extract information in a form that is different than you desired.

List within Lists

As was said before, the components of a list can be another list. Creating one is done in the same way we have already created a list.

```
listOfLists <- list(newList, "New Data")
listOfLists

[[1]]
[[1]]$LN
[1] 1 2 3 4 5 6 7 8 9 10

[[1]]$LC
[1] "a" "b" "c" "d" "e" "f"

[[1]]$LL
[1] TRUE FALSE TRUE FALSE

[[2]]
[1] "New Data"

str(listOfLists)

List of 2
 $ :List of 3
 ..$ LN: int [1:10] 1 2 3 4 5 6 7 8 9 10
 ..$ LC: chr [1:6] "a" "b" "c" "d" ...
 ..$ LL: logi [1:4] TRUE FALSE TRUE FALSE
 $ : chr "New Data"
```

Notice that there are two levels to this list. It is not that easy to see when displaying the list. However, looking at the `str()` of the list, we can see that `listOfLists` is a list made up of two components. The first component, listed after the first dollar sign, is a list with three components. The second component, listed after the second dollar sign, is a single piece of character data.

Vectors & Lists

We have been rather loose with the term vector. In *R*, vectors come in two flavors: atomic vectors and lists. Atomic vectors are those vectors that contain one type of data. Lists are vectors that can contain many types. We have a special word, lists, for the second type, it won't cause a problem referring to atomic vectors as 'vectors'.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Make a list called **Lists001**. It should have five components. None of the components should have names. Display **ONLY** the structure of **Lists001**. The content of each component is as follows:
 - a. `CharacterData001`
 - b. `Coerce05`
 - c. `MissingValues01`
 - d. `VectorsZero`
 - e. `PossiblyNAN`
2. Make a list called **Lists002**. It should have three components. The components should be named. Display **ONLY** the structure of **Lists002**. The content of each component is as follows:
 - a. `Chopper`: a vector with length 47 that starts at -23 and ends at -178. The values should be evenly spaced.
 - b. `Hera`: a vector of length 30 with all TRUE values.
 - c. `UpperBackwards`: a single string with the upper case alphabet recorded in reverse alphabetical order.
3. Make a list called **Lists003**. It should have two components. Only the first component should be named. Display **ONLY** the structure of **Lists003**. The content of each component is as follows:
 - a. `fruityFoods`: The matrix `fruitMatrix`
 - b. A character vector with 10 elements. All but the first are empty strings. The first value is missing.
4. Make a list called **Lists004**. It should have four components. The components should be named. Display **ONLY** the structure of **Lists004**. The content of each component is as follows:
 - a. `InsideList`: A list containing two unnamed components.
 - i. A single string of the character "a".
 - ii. A numerical vector with two values: 0, 1.
 - b. `InsideMatrix`: a matrix with 2 rows and three columns. All the values should be ones.
 - c. `InsideMissing`: A vector with length 10. All elements should be missing. It should have type character.

{MoreListssometime}

Chapter 17: Modifying Lists

Data From Lists

When extracting a multiple components from a list, the result can be **preserved** as a list, or can be possibly **simplified** into a ‘simpler’ data type. The choice of which method to use will depend on how the result will be used.

```
namedList

$LN
[1] 1 2 3 4 5 6 7 8 9 10

$LC
[1] "a" "b" "c" "d" "e" "f"

$LL
[1] TRUE FALSE TRUE FALSE
```

Preservation

To preserve the data as a list, use the name of the list followed by a single set of square brackets. Inside the square brackets, include a vector with the indices or the names of the desired components. A single sublist is extracted with just a numerical index or a sublist name.

```
namedList[1]
namedList["LN"]
```

The result of extracting a single sublist, using a single set of square brackets will be a list.

```
$LN
[1] 1 2 3 4 5 6 7 8 9 10

[1] "list"
```

Operations and functions that perform perfectly well on vectors, might not work on lists.

```
namedList[1] + 1

Error in namedList[1] + 1: non-numeric argument to binary operator
```

If desired, multiple components can be filtered out. This is done using a vector of indexes or a vector of names.

```
namedList[c(1, 3)]
namedList[c("LN", "LL")]

$LN
[1] 1 2 3 4 5 6 7 8 9 10

$LL
[1] TRUE FALSE TRUE FALSE
```

Simplification

Another way to extract the information from a list is to extract and attempt to simplify it. When *R* simplifies a list, if attempts to coerce the list into a ‘simpler’ form, from *R*s point of view. Simplification, of a single sublist, is attempted either by using double square brackets or a dollar sign followed by the component name.

```
namedList[[3]]
namedList$LL

[1] TRUE FALSE TRUE FALSE

typeof(namedList$LL)

[1] "logical"
```

In either case, for this example, the result is just a logical vector.

The `unlist()` function provides a way to collapse the components of a list into a single vector regardless of the number of components.

```
unlist(namedList[c(1, 3)])
```

| LN1 | LN2 | LN3 | LN4 | LN5 | LN6 | LN7 | LN8 | LN9 | LN10 | LL1 | LL2 | LL3 | LL4 |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|------|-----|-----|-----|-----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 1 | 0 | 1 | 0 |

Modifying List Components

Modifying list components can be done by assigning a new set of values to a particular component. This can be done either with names or numerical index.

```
namedList[2] <- TRUE
namedList$LC <- TRUE
namedList[[c(1, 1)]] <- 2
str(namedList)
```

```
List of 3
 $ LN: num [1:10] 2 2 3 4 5 6 7 8 9 10
 $ LC: logi TRUE
 $ LL: logi [1:4] TRUE FALSE TRUE FALSE
 - attr(*, "comment")= chr "This is a comment on a list."
```

Adding components

Adding Components can be done three ways. Two use the assignment operator. In either of these cases, you assign your new data to a component that doesn't already exist. Either refer to the next component index, or assign a new name using the dollar sign.

```
newDataIndex <- length(namedList)
namedList[newDataIndex + 1] <- "New Data"
namedList$ExtraData <- "Newer Data"
str(namedList)
```

```
List of 5
 $ LN      : num [1:10] 2 2 3 4 5 6 7 8 9 10
 $ LC      : logi TRUE
 $ LL      : logi [1:4] TRUE FALSE TRUE FALSE
 $         : chr "New Data"
 $ ExtraData: chr "Newer Data"
 - attr(*, "comment")= chr "This is a comment on a list."
```

When referring to the next component index, it is important to check the length of the list. Otherwise, several empty list components can be added.

The third method for adding a component to a list uses the `append()` function. It requires the list you have, and the components you want to add to it. This method requires that you assign your result to a variable, otherwise the change will not take place. The previous methods update the existing list.

```
namedList <- append(namedList, "Even More New Data")
str(namedList)
```

```
List of 6
 $ LN      : num [1:10] 2 2 3 4 5 6 7 8 9 10
 $ LC      : logi TRUE
 $ LL      : logi [1:4] TRUE FALSE TRUE FALSE
 $         : chr "New Data"
 $ ExtraData: chr "Newer Data"
 $         : chr "Even More New Data"
```

Removing Components

To remove a component, assign the value `NULL` to that component. This example will remove all but the first sublist.

```
namedList[2:6] <- NULL
str(namedList)
```

```
List of 1
 $ LN: num [1:10] 2 2 3 4 5 6 7 8 9 10
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Determine the structure of the following:
 - a. `ModifyLists001`
 - b. `ModifyLists002`
 - c. `ModifyLists003`
2. If possible, extract the indicated component(s) from each of the following lists. The result should remain a list. Display the structure of the result and not the result itself.
 - a. The first component of `ModifyLists001`
 - b. The last component of `ModifyLists002`
 - c. The third component of `ModifyLists003`
3. If possible, extract the indicated component(s) from each of the following lists. The result should remain a list.
 - a. The first two components of `ModifyLists001`
 - b. The last three components of `ModifyLists002`
 - c. The even components of `ModifyLists003`
4. If possible, extract the indicated component(s) from each of the following lists. The result should be a vector.
 - a. The first component of `ModifyLists001`:
 - b. The burger component of `ModifyLists002`
 - c. The graflex component of `ModifyLists003`
 - d. The first two components of `ModifyLists001`
 - e. The last five component of `ModifyLists002`
 - f. The even components of `ModifyLists003`
5. If possible, extract the indicated values(s) from each of the following lists.
 - a. The last two elements of the third component of `ModifyLists001`.
 - b. The first five **wild** elements of `ModifyLists002`.
 - c. The second through fourth element of **feet** in `ModifyLists003`. The result should be a character vector.
 - d. The elements of **feet** in `ModifyLists003` that are less than 10000 numerically. The result should be a character vector and not missing.
6. Remove the unnamed elements of `ModifyLists003`. Display the structure of the result. Do not modify the original list. Make a copy of it before you remove the elements.
7. Create a new list **CreatedList001** with three components. The components should be `ModifyLists001`, `ModifyLists002`, and `ModifyLists003`.
8. Create a new list **CreatedList002** with all components of the lists `ModifyLists001`, `ModifyLists002`, and `ModifyLists003`.

Chapter 18: DataFrames

When performing data analysis, most likely you will be using a **dataframe** to hold your data. A dataframe is a list where each component has the same length. To visualize a dataframe, think of a matrix whose columns can hold different types of data. The columns correspond to the components of a list. Generally, each column represents a single variable.

Creating a dataframe

Dataframes can be created many ways. Importing data is one way. But, the simplest way to make a dataframe is from a collection of vectors.

```
numbers <- 1:26
characters <- letters[1:26]
logicals <- rep(c(T, F), times = 13)
```

Using the `data.frame()` function, list the vectors that you want to include, separated by commas. Names can be included, if desired.

```
dat <- data.frame(numbers, charName = characters, 101:126, `2 name` = logicals)
# `2 name` is an incorrectly formatted name. The back-ticks 'allow' it.
head(dat, n = 4) # Change 6 to number of lines you want see.
```

```
  numbers charName X101.126 X2.name
1      1         a      101    TRUE
2      2         b      102   FALSE
3      3         c      103    TRUE
4      4         d      104   FALSE
```

Above, you can see the structure of a dataframe. You should notice, that the values in the *numbers* and *characters* vectors each have 26 elements. The dataframe as displayed only shows the first four rows. The `head()` function was used to display just the beginning of the dataframe. Entering *dat* alone, without using `head()`, would print all 26 lines. In this case, that is not much to look at, but consider a data frame with 10,000 rows. The output can be overwhelming. There is a `tail()` function. It displays the end of a dataframe. The `dim()` function will display the number of rows and columns.

If we continue to look at the structure of the dataframe, we will notice some behavior of `data.frame()`.

```
str(dat)

'data.frame': 26 obs. of 4 variables:
 $ numbers : int  1 2 3 4 5 6 7 8 9 10 ...
 $ charName: chr  "a" "b" "c" "d" ...
 $ X101.126: int  101 102 103 104 105 106 107 108 109 110 ...
 $ X2.name : logi  TRUE FALSE TRUE FALSE TRUE FALSE ...
```

First, `data.frame()` assigns a name to each component that is added. (Remember, these components are each columns.) The name for each column will come from either

1. the named variable (`numbers`)
2. a name given prior to naming the variable (`charName`)
3. some convention *R* uses for making up a name, if the first two types of names are not available.

Secondly, notice that `data.frame()` changed the name of a component. An *X* was appended to ``2 name`` because it starts with number. Spaces are replaced with a period. Basically, improperly formatted names are ‘fixed’. Improperly formatted names are common in imported data sets.

Finally, if we check the attributes of a dataframe, we can see that besides names, it has two other attributes namely *class* and *row.names*.

```
attributes(dat)

$names
[1] "numbers" "charName" "X101.126" "X2.name"

$class
[1] "data.frame"

$row.names
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26
```

Subsetting Dataframes

Elements of a dataframe are identified by both the row and column they are in. Using numerical vectors to select a single element, identify the row index and the column index. The “d” is in the fourth row and second column of the dataframe **dat**. To select “d”, type the name of the dataframe with a set of square brackets. Inside the brackets, enter the row index, a comma, and the column index.

```
dat[4, 2]
[1] "d"
```

If you want to restrict yourself to specific columns, at least two methods exist.

1. In the square brackets, enter a vector of column indexes or names. (This will return a dataframe.)
2. Replace the square brackets with a dollar sign, followed by the column name. (This will return a vector.)

```
dat[, c(1, 3)] # The leading comma is not needed when subsetting columns
dat[, c("numbers", "X101.126")]
dat$numbers
```

To replace a single value, a set of values, entire rows, or entire columns, use the assignment operator and reference the

Subsetting Entire Rows

Subsetting specific rows can be done in several different ways. After the vector name, and surrounded by square brackets

1. enter a numeric vector of desired rows, followed by a comma. (A vector of negative values will remove rows)
2. enter a logical vector followed by a comma. *TRUE* values will be kept.
3. enter a comparison, based on a column, followed by a comma. The comparison is usually based on values in a specific column.

As with vectors, the *subset()* function can also be used. The first argument is a dataframe, the second is the subset selection condition(s).

```
dat[c(1, 3), ]
dat[c(T, F, T, rep(F, times = 23)), ]
dat[dat$charName == "a" | dat$charName == "c", ]
subset(dat, dat$charName == "a" | dat$charName == "c")
```

Each of these filters out the same two rows.

| | numbers | charName | X101.126 | X2.name |
|---|---------|----------|----------|---------|
| 1 | 1 | a | 101 | TRUE |
| 3 | 3 | c | 103 | TRUE |

Missing Values

Determining which rows of a dataframe have missing values can be important. *complete.cases()* will return a logical vector where *FALSE* will indicate rows with missing values. *na.omit()* will return a dataframe with incomplete rows removed. However, columns of interest could be complete. In that case, removing rows is a bad idea.

Extending a dataframe

New columns can be added to a dataframe in several ways:

1. Use *data.frame()* listing the original dataframe with any new columns. Assign the result to a variable.
2. Using the dollar sign, add a new name to dataframe and assign the new data to it.
3. Use *cbind()* to bind a column to the existing data frame. Assign the result to a variable.

```
dat <- data.frame(dat, newNumbers1 = 1:26)
dat$newNumbers2 <- 11:36
dat <- cbind(dat, newNumbers3 = 21:46)
head(dat, n = 1)

  numbers charName X101.126 X2.name newNumbers1 newNumbers2 newNumbers3
1       1         a      101    TRUE           1          11           21
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Construct a single dataframe that has ALL of the following properties. The construction of each should not be done by brute force (listing out elements).
 - a. Name: `example.dataframe.1`,
 - b. Four variables with 10 observations each.
 - c. `variable.1` contains ten replications of the string 1. This should be caled `variable.1`.
 - d. `variable.2` contains numbers 1 - 10 in increasing order. This should be caled `variable.2`.
 - e. `variable.3` contains alternating logical values starting with `TRUE`. This should be caled `variable.3`
 - f. `variable.4` contains the first 10 colors from the `colors` function. This should be caled `variable.4`.
2. Construct a single dataframe that has ALL of the following properties. The construction of each should not be done by brute force (listing out elements).
 - a. Name: `example.dataframe.2`,
 - b. Three variables with 78 observations each.
 - c. `variable.1` contains the letters of the lowercase alphabet each repeated 3 times consecutively.
 - d. `variable.2` contains numbers 1 - 78 in decreasing order.
 - e. `variable.3` contains the last 78 colors in the `colors()` function. Starting with the last in the list alphabetically.
3. How many rows and columns do these dataframes have? This should be done with a function, not by inspection.
 - a. `ChickWeight`
 - b. `Formaldehyde`
 - c. `Loblolly`
4. Extract the following from the indicated dataframe: (This should not be done by inspection. The result should be a `data.frame`.)
 - a. `Loblolly`: The age variable
 - b. `Formaldehyde`: The optical density variable
 - c. `ChickWeight`: The weight variable
5. Extract the following from the indicated dataframe: (The result should be a vector.)
 - a. `Loblolly`: The age of all `Loblolly` pine trees with a height over 55 ft.
 - b. `Formaldehyde`: The optical density when the Carbohydrate is greater than 0.4 ml
 - c. `ChickWeight`: The weight of chicks at 4 days old.
6. Extract the following from the indicated dataframe: (This should not be done by inspection. The result should be a `dataframe`)
 - d. `Loblolly`: Rows with heights over 60 ft
 - e. `Formaldehyde`: Rows with carbs over 0.5 ml
 - f. `ChickWeight`: Rows at age 0 days.

Questions. Questions. Questions. I need more questions

Chapter 19: Tibbles

A **tibble** is a modified version of a **dataframe**²². The *tibble* package adds the necessary functionality to R.

```
library(tibble)
```

Creating a tibble

As with dataframes, tibbles can be created in different ways. They are the result of importing data. They are the result of using the *tibble()* function in a manner analogous to using *data.frame()* to bind vectors together. For comparison, a tibble and a dataframe are constructed from the same set of vectors. As with a dataframe, list the vectors, possibly with names, separated by commas²³.

```
numbers <- 1:26
characters <- letters[1:26]
logics <- rep(c(T, F), length = 26)
# Building a dataframe
dat <- data.frame(numbers, charName = characters, 101:126, `2 name` = logics)
# Building a tibble.
tib <- tibble(numbers, charName = characters, 101:126, `2 name` = logics)
```

A tibble can also be built by directly typing in the values row by row. This is done with the *tribble()* function. To do so, list the column names separated by a comma in the first row of entered data. Immediately precede each name with a tilde. One subsequent line enters the data row by row separating values with a comma.

```
tribble.example <- tribble(~var.1, ~var.2, 1, "q", 2, "w", 30, "e")
tribble.example

# A tibble: 3 × 2
  var.1 var.2
  <dbl> <chr>
1     1 q
2     2 w
3    30 e
```

A tibble can also be made directly from a dataframe, vector, or matrix. The *as_tibble()* function will attempt to coerce these into a tibble. Going the other direction, the *as.data.frame()* function can be used to coerce a tibble into a dataframe²⁴.

Differences between tibble and dataframe

For comparison, the tibble **tib** and dataframe **dat** are displayed below:

```
tib

# A tibble: 26 × 4
  numbers charName `101:126` `2 name`
  <int> <chr>      <int> <lgl>
1     1 a          101 TRUE
2     2 b          102 FALSE
3     3 c          103 TRUE
4     4 d          104 FALSE
5     5 e          105 TRUE
6     6 f          106 FALSE
7     7 g          107 TRUE
8     8 h          108 FALSE
9     9 i          109 TRUE
10    10 j          110 FALSE
# [i] 16 more rows

# `2 name` is an incorrectly formatted name. The back-ticks 'allow' it.
head(dat, n = 4)

  numbers charName X101.126 X2.name
1     1     a      101     TRUE
2     2     b      102     FALSE
3     3     c      103     TRUE
4     4     d      104     FALSE
```

²² Most of this information is drawn from `_vignette("tibble")`.

²³ For tibbles, *tibble()* will only recycle vectors of length one.

²⁴ Older functions may not recognize tibbles

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Construct a single tibble that has ALL of the following properties. The construction of each should not be done by brute force (listing out elements).
 - a. Name: `example.tibble.1`,
 - b. Four variables with 10 observations each.
 - c. `variable.1` contains ten replications of the string 1. This should be caled `variable.1`.
 - d. `variable.2` contains numbers 1 - 10 in increasing order. This should be caled `variable.2`.
 - e. `variable.3` contains alternating logical values starting with `TRUE`. This should be caled `variable.3`
 - f. `variable.4` contains the first 10 colors from the `colors` function. This should be caled `variable.4`.
2. Construct a single tibble that has ALL of the following properties. The construction of each should not be done by brute force (listing out elements).
 - a. Name: `example.tibble.2`,
 - b. Three variables with 78 observations each.
 - c. `variable.1` contains the letters of the lowercase alphabet each repeated 3 times consecutively.
 - d. `variable.2` contains numbers 1 - 78 in decreasing order.
 - e. `variable.3` contains the last 78 colors in the `colors()` function. Starting with the last in the list alphabetically.
3. How many rows and columns do these tibbles have? This should be done with a function, not by inspection.
 - a. `ChickWeight.tib`
 - b. `Formaldehyde.tib`
 - c. `Loblolly.tib`
4. Extract the following from the indicated tibble: (This should not be done by inspection. The result should be a tibble.)
 - a. `Loblolly.tib`: The age variable
 - b. `Formaldehyde.tib`: The optical density variable
 - c. `ChickWeight.tib`: The weight variable
5. Extract the following from the indicated tibble: (The result should be a vector.)
 - a. `Loblolly.tib`: The age of all Loblolly pine trees with a height over 55 ft.
 - b. `Formaldehyde.tib`: The optical density when the Carbohydrate is greater than 0.4 ml
 - c. `ChickWeight.tib`: The weight of chicks at 4 days old.
6. Extract the following from the indicated tibble: (This should not be done by inspection. The result should be a tibble)
 - d. `Loblolly.tib`: Rows with heights over 60 ft
 - e. `Formaldehyde.tib`: Rows with carbs over 0.5 ml
 - f. `ChickWeight.tib`: Rows at age 0 days.
7. Create a tibble **alpha.tib** where the first component **forwardAlpha** is a vector of all the lowercase letters in the English alphabet. The second component **backwardAlpha** should be the Uppercase letters of the English alphabet in reverse alphabetical order. The third component **bothAlpha** should be a function of the first two components where each element is a character string of length two consisting of the corresponding elements of the first two components.

Questions. Questions. Questions. I need more questions

Chapter 20: Functions

When coding, you may find that with minor modifications, blocks of code are being reused. This will be most apparent when a chunk of code is copied, pasted, and a few values are changed. Once this occurs, a function should be made out of this code. A **function** is a named set of commands. This has at least two advantages:

1. Code is easier to read. Repetitive sets of commands are replaced with a function call and its arguments.
2. Code is easier to correct. Without making a function, an error in a repeated chunk of code must be corrected in each replication. With a function, the correction only needs to be made once, where the function is written.

Example

In previous sections, many examples of functions have been encountered. Some of these include: *any()*, *print()*, *c()*.

With the exception of the last, these functions come built into the basic *R* installation. The last is added using the *dplyr* package.

Even if we can not see them, each of these follows a set of commands, and outputs a result.

Basic Form

This section will deal with creating new functions²⁷. When creating a new functions, it will have this basic form:

```
functionName <- function( arguments = defaults ){
  steps to be executed
  return( optional output to be returned )
}
```

It starts with *functionName*²⁸. This is the name of the function. This name will be used whenever you want to execute the steps contained within the vector. The function name is followed by the assignment operator, and then the **function()** command²⁹. This command indicates to *R* that a function is being built. Within the round parentheses, list any arguments for the function.

Arguments

Arguments are variables used to customize how a function works. Most, but not all, functions that are written will require some arguments. If needed, some of the arguments can be assigned default values. These default values will be used within the function unless otherwise specified. A default is made by setting the argument equal to the default value

argument = default value

Body of a Function

Following the round parentheses, that contain the arguments, will be a set of curly braces. Within the curly braces will be the code the function is meant to execute. The last line inside the curly braces is a *return()* statement. The *return()* statement will specifically indicate what the output will be for the function.

As an example, the function *close()* is given below.

```
# Determines if two values x and y are within tol of each other.
close <- function(x, y, tol = 1e-04) {
  abs.diff <- abs(x - y) #step 1
  less.than.tol <- abs.diff < tol #step #2
  return(less.than.tol)
}
```

As the name, and comment, indicate, *close()* determines if two values are close to each other.

Step #1 Determine the absolute value of the difference of *x* and *y* and assign this value to a variable *abs.diff*.

Step #2 Determine if *abs.diff* is less than the variable *tol*. Assign the logical result to the variable *less.than.tol*. The last line returns the value for *less.than.tol*.

²⁷ New functions need to be created in the code BEFORE it is used in the code.

²⁸ Anonymous functions are functions where no *functionName* is assigned.

²⁹ **function()** can be abbreviated with **\()**

The description of the function leaves x , y and tol unspecified. These are the arguments. The x and y arguments will specify the values being compared. The tolerance tol specifies what it means for two values to be close.

In `close()`, tol 's default is set to 0.0001. Meaning that unless tol is explicitly changed, 0.0001 will be used.

```
first.Vector <- 1.01
second.Vector <- 1.02

close(x = first.Vector, y = second.Vector)

[1] FALSE

close(x = first.Vector, y = second.Vector, tol = 0.1)

[1] TRUE
```

Calling a Function

When using a function, the arguments can be listed in several ways.

1. When using argument names, the arguments can be listed in any order, as long as they are separated by a comma.
2. When not using argument names, the values need to be entered in the same order as the argument.

These would run through the same computations:

```
close(first.Vector, second.Vector)
close(x = first.Vector, y = second.Vector, 0.1)
```

However, this would not.

```
close(first.Vector, 0.1, second.Vector)
```

In this case, the x argument takes values from *first.Vector*, the y argument is set to 0.1, and tol is set to *second.Vector*.

Multiple Outputs

Functions produce only a single result. The result can be any data type: a vector, a list, a dataframe, or nothing at all.

A modified `close()` function, returns a dataframe containing the absolute difference, the tolerance, and the determination.

```
# Determines if two values x and y are within tol of each other.
close.2 <- function(x, y, tol = 1e-04) {
  abs.diff <- abs(x - y) #step 1
  less.than.tol <- abs.diff < tol #step #2
  output <- data.frame(difference = abs.diff, tolerance = tol, closeTF = less.than.tol)
  return(output)
}

close.2(first.Vector, second.Vector)

  difference tolerance closeTF
1      0.01    1e-04    FALSE
```

Functional Programming

R is a functional programming language. This means that a functions argument can be another function. For example, the function `fun()` has two arguments fn and x .

```
fun <- function(fn, x) {
  computed <- fn(x)
  return(computed)
}
```

Examining the code, it can be seen that `fun()` applies the function $fn()$ to the value x . The example below shows `fun()` applying `mean()` and `min()` to the vector `c(2,1,3)`.

```
fun(fn = mean, x = c(2, 1, 3))
fun(fn = min, x = c(2, 1, 3))

[1] 2
[1] 1
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Make a function **squarer** that takes a single numerical argument x and returns the square of the value of x . Assume that only numerical values for x will be supplied.
2. Make a function **FtoC** that takes a single numerical argument $Fahr$ and returns the Celsius equivalent. Assume that only numerical values for $Fahr$ will be supplied.
3. Make a function **rootOf100** that takes a single numerical argument n and returns the n -th root of 100. Assume that only numerical values for n will be supplied. The default for n should be one.
4. Make a function **CtoF** that takes a single numerical argument $Celc$ and returns the Celsius equivalent. Assume that only numerical values for $Celc$ will be supplied. The default for $Celc$ should be 0.
5. Make a function **Bigger.Logical** that takes two numerical arguments x and y . Assume that only numerical values for x and y will be supplied. It should return a logical value of TRUE if x is larger than y . It should return a FALSE otherwise.
6. Make a function **Bigger.Numeric** that takes two numerical arguments x and y . Assume that only numerical values for x and y will be supplied. It should return a numerical value of 1 if x is larger than y . It should return a zero otherwise.
7. Make a function **Smaller** that takes two numerical arguments x and y . Assume that only numerical values for x and y will be supplied. It should return the smaller of the values x and y .
8. Make a function **totalOfTen** that a single numerical argument x . Assume that only numerical values of length ten will be supplied. **totalOfTen** should return a vector with a single element in it. The value of that single element should be the total of all ten values in x . The default value for x should be a vector of ten zeros.

Questions. Questions. Questions. I need more questions

Chapter 21: Controls - Decisions

If

An *if* statement takes the form

```
if( condition ){
  Steps to execute IF the condition is TRUE
}
```

The statement will check a condition. The condition, contained in round parentheses, should evaluate to a single logical value. If that logical value is *TRUE*, the code inside the curly braces will be executed. If that logical value is *FALSE*, the code inside the curly braces will be skipped.

Looking over these two snippets of code, both *if* statements determine if all values in a given vector are large than zero. If they are, *R* will print out a message. Otherwise, *R* will do nothing.

```
numbers = c(1, 8)

if (all(numbers > 0)) {
  print("The numbers are all positive.")
}

[1] "The numbers are all positive."

numbers = c(-1, 1)

if (all(numbers > 0)) {
  print("The numbers are all positive.")
}
```

As can be seen, only the first snippet produces a result.

If...Else...

An *if...else...* statement expands on the *if* statement. It allows alternatives outcomes to be produced. An *if...else...* strings along a series of *if* statements. The first time *R* finds an *if* statement with a *TRUE* condition, it will execute the code in the accompanying curly braces. Once *R* executes that particular set of commands, it will not check anymore conditions in the *if...else...*, and will skip paste any more accompanying code.

```
numbers = c(-1, -8)

if (all(numbers > 0)) {
  print("The numbers are all positive.")
} else if (all(numbers < 0)) {
  print("The numbers are all negative.")
} else if (all(numbers < 0)) {
  # An intentional mistake.
  print("The numbers are all zero.")
} else {
  print("The numbers are a mess.")
}

[1] "The numbers are all negative."
```

In this snippet, *R* determines if all the numbers are negative. Since the first condition fails, *R* checks the second. It is *TRUE*, therefore *R* prints “The numbers are all negative.” At this point *R* does not check any further conditions. This can be seen by looking at the third condition. It is also a *TRUE* statement, but the accompanying statement is not printed.

For both an *if* statement, and an *if...else...*, multiple lines of code can be included for each condition. A single line was used in the example for brevity.

ifelse

The `ifelse()` function is a vectorized version of a single `if...else...` statement. The syntax for `ifelse()` is

`ifelse(logicalcondition,result.T, result.F).`

`ifelse()` will evaluate the condition, and return `result.T` if the condition is `TRUE`, otherwise it will return `result.F`.

```
x <- c(-1, 0, 3)
ifelse(x > 0, "Positive", "Not Positive")

[1] "Not Positive" "Not Positive" "Positive"
```

Several conditions can be evaluated using `ifelse()` by nesting an `ifelse()` inside of itself.

`ifelse(condition1,result.T1,ifelse(condition2,result.T2,result.F))`

```
x <- c(-1, 0, 3)
ifelse(x > 0, "Positive", ifelse(x < 0, "Negative", "Zero"))

[1] "Negative" "Zero"      "Positive"
```

However, this can be very hard to read. Readability will also be very difficult if the results are several lines of code.

switch

`switch()` can be used to select from multiple results. The basic form is

`switch(EXPR,result1,result2,...).`

`switch()` uses its first argument `EXPR` to select the result from the successive arguments. The results are listed in one of two distinct ways depending whether the first argument evaluates to a character string or a number.

In this first example, `EXPR` evaluates to a number. The result is determined by counting through the listed results.

```
x <- 2
switch(EXPR = x, "a", "b", "c", 100) #The second result is 'b'.

[1] "b"

x <- 4
switch(EXPR = x, "a", "b", "c", 100) #The fourth result is 100.

[1] 100
```

Take note that the results are a list of values, but they are not collected into a `list()` or a vector.

In this second example, `EXPR` evaluates to a string. For this to evaluate correctly, each result, with one exception, needs to be given a name. `R` will select the result whose name matches `EXPR`.

```
x <- "b"
switch(EXPR = x, a = 1, b = 2, c = "c", 100) # 'b' matches with the result labeled b.

[1] 2

x <- "e"
switch(EXPR = x, a = 1, b = 2, c = "c", 100) #EXPR = 'e' has no match. The unnamed result is used..

[1] 100
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Write an if-else statement that will print the string, "Controls001 has a number eight", if the vector Controls001 contains the number eight, otherwise it will print the string "Controls001 does not have a number eight"
2. Write an if-else statement that will print the string, "Controls001 has a negative number", if the vector Controls001 contains a negative number, otherwise it will print the string "Controls001 does not have a negative number"
3. Write an if-else statement that will print the string:
 - a. "Controls001 has all positive numbers", if the vector Controls001 contains all positive numbers,
 - b. "Controls001 has all negative numbers", if the vector Controls001 contains all negative numbers numbers,
 - c. "Controls001 has all zeros", if the vector Controls001 contains all zeros,
 - d. "Controls001 is a mess" for anything else
4. Write an if-else statement that will print the string:
 - a. "The first element of Controls002 is the smallest", if this is true
 - b. "The second element of Controls002 is the smallest", if this is true
 - c. "The third element of Controls002 is the smallest", if this is true
 - d. "The fourth element of Controls002 is the smallest", if this is true
5. Create a function **Bigger** that takes two arguments: x and y . It should return the larger of the two values. If they are the same, it should return x . Assume that x and y will only have vectors of length one passed to each of them. This should make use of an if/else statement.
6. Create a function **BiggerSquaredOrHalf** that takes two arguments: x and y . It should return the the square of x if x is greater. Otherwise, it should return one-half of y . Assume that x and y will only have vectors of length one passed to each of them. This should make use of an if/else statement.
7. Add two variables to the **LifeCycleSavingsUpdate** dataframe: **age.ratio** and **broken.up**
 - a. **age.ratio** should be the ratio of the **pop15** variable to the **pop75** variable.
 - b. **broken.up** should take one of three values: "UP", "DOWN", or "OVER". Use "DOWN" if **age.ratio** is less than 10 and **ddpi** is less than 3. Use "UP" if **age.ratio** is greater than 20 and **ddpi** is greater than 4. For everything else, use "OVER".
8. Using the **oldFaithful** dataframe, print out one of the following strings:
 - a. "Old faithful has XXX more eruptions that last at least four minutes than eruptions that last less than three minutes."
 - b. "Old faithful has XXX more eruptions that last less than four minutes than eruptions that last at least than three minutes."
 - c. "Old faithful has the same number of eruptions that last less than four minutes as eruptions that last at least than three minutes. There are XXX each."
 - d. XXX should be replaced by the positive difference in the number of eruptions that last at least four minutes versus those that last less than four minutes. If they are the same, XXX should be replaced by zero.
9. Make a function `tempChange()` that meets the following requirements:
 - a. It takes two arguments: *givenTemp* and *to.C*.
 - b. *givenTemp* should have no default. *to.C* should default to `TRUE`.

- c. If *to.C* is *TRUE*, then *tempChange*() should convert *givenTemp* from a Fahrenheit temperature to its equivalent temperature in Celcius.
- d. If *to.C* is *FALSE*, then *tempChange*() should convert *givenTemp* from a Celcius temperature to its equivalent temperature in Fahrenheit.
- e. *tempChange*() should return a list with three named components: *givenTemp*, *to.C*, and *convertedTemp*. *convertedTemp* is the converted temperature. Demonstrate the usage of your function by applying it to the **changeTheseTempsVector** using *to.C* set to both options.

Questions. Questions. Questions. I need more questions

Chapter 22: Controls - Loops

A **loop** is a piece of code that will be repeat the commands contained inside of it. Writing a loop can make your code shorter, since the code is not repeated, and it will make it easier to read. Situations where a loop will be useful can be broken down into two cases:

1. The number of repetitions/iterations can be determined before the looping process begins.
2. The number of repetitions/iterations can not be determined until after the loop begins.

(In reality, both cases can be handled by either of the functions that will be outlined. But, it is easier to think of them as separate cases)

for() Loop

A *for*() loop is a looping method used when the number of repetitions/iteration can be determined before the looping procedure begins. As an example, a *for*() loop will be constructed to print each element in a given vector. It will be shown once without using a *for*() loop, and once with.

```
print(letters[1])
print(letters[2])
print(letters[3])
print(letters[4])
print(letters[5])

for (index in 1:5) {
  print(letters[index])
}

[1] "a"
[1] "b"
[1] "c"
[1] "d"
[1] "e"
```

Both pieces of code produce the same output: the first five characters in *letters*. The second is much shorter. In fact, if the goal was to print all 26 characters in *letters*, the first method would require 26 lines of code. The second method would need to change the 5 to a 26.

The form of a *for*() loop,

```
givenVector <- c(element1, element2, ..., element.Final)

# The for() Loop begins here.
for(eachElement in givenVector){
  # perform the given steps
}
```

makes it relatively straight forward to conceptualize the process.

For each element in a given vector, complete the given steps.

The syntax can be hard to conceptualize. It begins with the *for*() command. Inside the round parentheses are three objects, in order:

1. an indexing variable
2. the word 'in'
3. a given vector of any type.

The word 'in', in a certain sense, acts as an assignment operator. For each pass through the loop, it assigns the next value in the vector to the indexing variable. Since the given vector has a finite number of elements, the loop will only run a finite number of times.

Indexing Variable

The indexing variable does not need to be used explicitly in the steps contained in a *for*() loop, but it can be. When it is not used, it acts to count the repetitions. When it is used, in addition to counting the repetitions, it also customizes the code for each repetition.

Skipping Iterations - next/break

Under certain conditions, skipping a single repetition, or all remaining repetitions, will be the proper course of action. To do so, the body of the loop should include an *if*() statement whose condition, when *TRUE*, tells *R*, to advance the indexing variable to the next value, or leave the loop altogether. The command following the *if*() statement will be *next*, or *break*, respectively.

```
# The for() Loop begins here.
for(eachElement in givenVector){
  if( skipping.condition ){
    next
  }
  # perform the given steps, if the skipping condition is FALSE
}
```

while() Loop

A *while*() loop is a looping method used when the number of repetitions can't be determined before the looping procedure begins. The basic form is as follows:

```
while( given.Condition == TRUE){
  # perform the given steps
  # update check.Value
}
```

A *while*() loop will execute a given set of commands as long as a given condition is *TRUE*. Once the given condition becomes *FALSE*, *R* will skip past the code in the loop and continue with any subsequent code.

As an example, the *sample*() function will select a random value from 1 to 100. A *while*() loop will print a sentence about the random value, and then select a new one. As long as the random value is less than 90, *R* will continue to print and select new values. Since each successive number is unknown, it can not be determined ahead of time how many values will be printed.

```
randomValue <- sample(x = 1:100, size = 1)
while (randomValue < 90) {
  sentence <- paste("The current random value is ", randomValue, ".", sep = "")
  print(sentence)
  randomValue <- sample(x = 1:100, size = 1)
}

[1] "The current random value is 42."
[1] "The current random value is 83."
[1] "The current random value is 31."

randomValue

[1] 92
```

The last value that is printed is a 92. But, it is not part of sentence as the other results are. This is because once the 92 was selected, the condition was no longer *TRUE*. When the condition is *FALSE*, *R* dropped out of the loop, and executed the command that displayed just the number 92.

Care should be taken when using a *while*() loop. If it is not written properly, it is possible for the loop to continue infinitely. This code, if run, will never stop. Replacing *randomValue* < 90 with *randomValue* < 1000 will give a condition that will never be *FALSE*. Theoretically, a loop with the second condition will run forever.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

1. Use a `for()` loop to add together the values in the vector `Loops001`.
2. Use a `for()` loop to add together the values in the vector `Loops001`. The result should be a vector with the same length as `Loops001`. The i^{th} value of the result should be the total of the first i values in `Loops001`.
3. Using a `for()` loop, create a vector with 40 elements. The first element should be -4. The second should be 3. Subsequent elements should be the total of the previous two elements.
4. Using a `for()` loop, create a vector with 40 elements. The first three elements should all be the number one. Subsequent elements should be the total of the previous three elements.
5. Using a `for()` loop, create a vector with 40 elements. The first element should be the number one. Subsequent elements should be the negative of two times the previous.
6. Use a `for()` loop, `LoopsMascots`, `LoopsPunc`, and `LoopsPhrase` to recreate these vectors of length one.

```
[1] "Once upon a time, my mascot was a Pink Turtle. I was five."
[1] "Once upon a time, my mascot was a Blue Devil! mehh."
[1] "Once upon a time, my mascot was a Tiger. ROAR!!!"
[1] "Once upon a time, my mascot was a Golden Eagle. Wow, I can fly."
[1] "Once upon a time, my mascot was a Fuzzy Orange? Yup, a Fuzzy Orange. "
[1] "Once upon a time, my mascot was a Big Red Bear. Really, there is no mascot."
```

```
[1] "Once upon a time, my mascot was a Pink Turtle. I was five." "Once upon a time, my mascot was a Blue Devil
! mehh." "Once upon a time, my mascot was a Tiger. ROAR!!!"
[4] "Once upon a time, my mascot was a Golden Eagle. Wow, I can fly." "Once upon a time, my mascot was a Fuzzy Oran
ge? Yup, a Fuzzy Orange. " "Once upon a time, my mascot was a Big Red Bear. Really, there is no mascot."
```

7. Use a `for()` loop, `LoopsMascots`, `LoopsPunc`, and `LoopsPhrase` to recreate these vectors of length one.

```
[1] "Once upon a time, my mascot was a Pink Turtle. I was five."
[1] "Once upon a time, my mascot was a Blue Devil! mehh."
[1] "Once upon a time, my mascot was a Tiger. ROAR!!!"
[1] "Once upon a time, my mascot was a Golden Eagle. Wow, I can fly."
[1] "Once upon a time, my mascot was a Fuzzy Orange? Yup, a Fuzzy Orange. "
[1] "Once upon a time, my mascot was a Big Red Bear. Really, there is no mascot."
```

```
[1] "Once upon a time, my mascot was a Pink Turtle. I was five." "Once upon a time, my mascot was a Blue Devil
! mehh." "Once upon a time, my mascot was a Tiger. ROAR!!!"
[4] "Once upon a time, my mascot was a Golden Eagle. Wow, I can fly." "Once upon a time, my mascot was a Fuzzy Oran
ge? Yup, a Fuzzy Orange. " "Once upon a time, my mascot was a Big Red Bear. Really, there is no mascot."
```


Questions. Questions. Questions. I need more questions


Chapter 23: Functions - Set Operations

A vector is a collection of elements. Usually, the order of the elements in the vector matters. We can state what the first, second, and third (and so on) elements are. However, in some instances, we are only interested in knowing what elements are included in a vector, or a collection of vectors. In this case, we the vectors can be thought of as sets.

In our example, we will define three character vectors.

```
x <- c("A", "B", "C", "D", "A", "A", "A")
y <- c("C", "D", "E", "F", "F", "E")
```

Here we have a visual illustration of these vectors when thought of as sets.




Union

The union of two sets is all elements that occur in at least one of the two sets. Looking at the union of vector x and y , the union is all elements except for "G".

```
union(x, y) #elements in at least one of x or y
```

```
[1] "A" "B" "C" "D" "E" "F"
```




Notice that the union of these two sets produces a vector that contains all the elements of vector x listed first. That is, with the exception of the final three "A"s in vector x . Then, elements from vector y are listed in the order that they appear in y , provided they are not already listed.

Intersection

The intersection of two sets is all elements that occur in both two sets. Looking at the intersection of vector x and y , the intersection is "C" and "D".

```
intersect(x, y) #elements in both x and y
```

```
[1] "C" "D"
```




It will be the case that the intersection of two sets will always be included in the union.

Set Difference

The set difference is like subtraction for sets. Take the first set, and remove from it all elements that are in the second. The remaining elements are the set difference. As with subtraction, the order of the subtraction will matter. In this case, removing all element of x that fall into y leaves a set difference of “E” and “F”.

```
setdiff(y, x) #elements in y but not in x
```

```
[1] "E" "F"
```



Switching the order around, the set difference would be “A” and “B”, if y was subtracted from x .

Set Equality

Set equal can be tricky. It checks for the existence of each element of one vector in the other, and vice versa. It is not checking that there are the same number of each, or that they appear in the same order. It only checks for existence.

The elements in x and y are not the same.

```
setequal(x, y)
```

```
[1] FALSE
```

Here, a new vector z is created. It is shorter than x but contains the same elements. Therefore, as sets, the two vectors are equal.

```
z <- c("A", "D", "B", "C")
```

```
setequal(x, z)
```

```
[1] TRUE
```

Recall, to determine if two vectors have the same elements, in the same order, use the identical function.

```
identical(x, z)
```

```
[1] FALSE
```

Is this %in% in the set?

The `%in%` operator determines if elements of one vector are *in* another. It will return a logical value for each element in vector listed on the right hand side of `%in%`.

```
"A" %in% x #is this element in this vector
```

```
[1] TRUE
```

```
"A" %in% y
```

```
[1] FALSE
```

```
c("A", "E", "A") %in% x
```

```
[1] TRUE FALSE TRUE
```

Along with a few other commands, the `%in%` operator can be used in a similar manner as `setequal()`.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

Questions. Questions. Questions. I need more questions

Chapter 24: Functions - Mathematical Operations

Sums & Products

: The `sum()` function totals the values given to it. The basic structure for using the function is

$$\text{sum}(x, na.rm = FALSE)$$

where x is a vector of numeric or logical values to be summed. The second argument `na.rm =` determines whether missing values should be removed or not. By default, they are included in the computation. When missing values are included, the result will be a missing value. If there are no missing values, the argument can be ignored.

```
x <- c(1, 5, 3, 2, NA, 4, NA)
sum(x, na.rm = FALSE)

[1] NA

sum(x, na.rm = TRUE)

[1] 15
```

The sum over multiple vectors can be taken. To do so, list all desired vectors in the `sum()` function, separating them with commas.

```
y <- c(0, 1)
sum(x, y, na.rm = TRUE)

[1] 16
```

Logical values can also be summed. Numerically, `TRUE` equals 1 and `FALSE` equals 0. If `TRUE` represents the occurrence of an event, summing a logical vector determines the total number of occurrences.

```
logical.x <- c(TRUE, TRUE, FALSE, TRUE)
sum(logical.x)

[1] 3
```

Cumulative Sums

The `cumsum()` function computes the cumulative sums of a vector. `cumsum()` returns a vector with the same length as the original vector. The first output element is the first element of the original vector. The second element outputted is the sum of the first two elements of the original vector. The third element outputted is the sum of the first three elements of the original vector. This repeats for each element of the original vector. `cumsum()` does not have an argument that deals with missing values. Once a missing value appears, the remaining cumulative sums are missing.

```
cumsum(x)

[1] 1 6 9 11 NA NA NA
```

Differences:

The `diff()` function takes different of vector elements whose indices are a fixed distance apart. This distance is called **lag**.

```
x <- c(1, 22, 333, 44, 5, 66, NA)
diff(x, lag = 2)

[1] 332 22 -328 22 NA
```

With a lag of 2, the difference is taken of the the first and third elements ($\text{lag} = 3 - 1 = 2$), the second and fourth elements ($\text{lag} = 4 - 2 = 2$), and so on. The vector of differences will be shorter than the one original vector.

Products:

The `prod()` function multiplies the values given to it. The basic structure for using the function is

$$\text{prod}(x, na.rm = FALSE)$$

where x is the vector of numeric or logical values to be summed. The second argument `na.rm =` determines whether missing values should be removed or not. By default, they are included in the computation. When missing values are included, the result will be a missing value. If there are no missing values, the argument can be ignored.

```
x <- c(1, 5, 3, 2, NA, 4, NA)
prod(x, na.rm = FALSE)

[1] NA

prod(x, na.rm = TRUE)

[1] 120
```

The product can be taken over more than one vector. To do so, list all desired vectors in the `prod()` function, separating them with commas.

```
y <- c(0, 1)
prod(x, y, na.rm = TRUE)

[1] 0
```

As with addition, Logical values can also be multiplied. If `TRUE` represents the occurrence of an event, multiplying all elements of a logical vector determines whether the event always occurred or if it failed to occur at least once.

```
logical.x <- c(TRUE, TRUE, FALSE, TRUE)
prod(logical.x)

[1] 0
```

Cumulative Products:

The `cumprod()` function computes the cumulative products of a vector. `cumprod()` returns a vector with the same length as the original vector. The first output element is the first element of the original vector. The second output element is the product of the first two elements of the original vector. The third output element is the product of the first three elements of the original vector. This repeats for each element of the original vector. `cumprod()` does not have an argument that deals with missing values. Once a missing value appears, the remaining cumulative sums are missing.

```
cumprod(x)

[1] 1 5 15 30 NA NA NA
```

Factorials:

The `factorial()` function takes the product of the first n positive integers.

```
factorial(5)

[1] 120

prod(1:5)

[1] 120
```

Binomial Coefficient:

Built from the factorials is the `choose()` function. It computes binomial coefficients, namely

$$\binom{n}{m} = \frac{n!}{m!(n-m)!}$$

where n, m are both non-negative integers with n greater than or equal to m .

```
choose(5, 2)

[1] 10

numerator <- factorial(5)
denominator <- factorial(2) * factorial(5 - 2)
numerator/denominator

[1] 10
```

Maximums

The `max()` function finds the maximum value within the vector given to it. The basic structure for using the function is

$$\text{max}(x, na.rm = FALSE)$$

where x is a vector of numeric or logical values. The second argument `na.rm` = determines whether missing values should be removed or not. By default, they are included in the computation. When missing values are included, the result will be a missing value. If there are no missing values, the argument can be ignored.

```
x <- c(30, NA, 12, 100, 43, 67, 100)
max(x, na.rm = FALSE)

[1] NA

max(x, na.rm = TRUE)

[1] 100
```

The maximum can be taken over more than one vector. To do so, list all desired vectors in the `max()` function, separating them with commas.

```
y <- c(0, 1000, 15, 20, 11, 0, 0)
max(x, y, na.rm = TRUE)

[1] 1000
```

The maximum of a logical vector can be found. *R* coerces (changes) *TRUE* to a value of 1, and *FALSE* as a 0. If *TRUE* represents the occurrence of an event. If the maximum of a vector equals 1, this implies that the event occurred at least once.

```
logical.x <- c(TRUE, TRUE, FALSE, TRUE)
max(logical.x)

[1] 1
```

Parallel Maximum :

The `pmax()` function is used to determine parallel maximum over a collection of vectors. `pmax()` will return a vector that contains the maximum of the first elements, the maximum of the second elements, and so on. The basic structure for the pairwise maximum is

$$\text{pmax}(\dots, na.rm = FALSE)$$

³⁰ where ... are the given vectors separated by a comma. The argument `na.rm` = determines whether missing values should be ignored or not. By default, they are not ignored in the computation. When missing values are included, the result will be a missing value. If there are no missing values, the argument can be ignored.

```
# This example only uses
x <- c(30, NA, 12, 100, 43, 67, 100)
y <- c(0, 1000, NA, 20, 11, 0, 12)
z <- 17
pmax(x, y, z, na.rm = FALSE)

[1] 30 NA NA 100 43 67 100

pmax(x, y, z, na.rm = TRUE)

[1] 30 1000 17 100 43 67 100
```

Notice that when `na.rm = TRUE`, the missing value is ignored. If the vectors represents by ... do not have the same length, the shorter will be recycled. If more than two vectors are used, list them, separating them with commas. This will return the parallel maximum.

³⁰ The 'p' in `pmax()` actually stands for parallel.

Index of Maximum:

The *which.max()* determines the index of the *first* occurrence of the maximum of a vector.

```
which.max(x)
```

```
[1] 4
```

Minimums:

The *min()* function finds the minimum value within the vector given to it. The basic structure for using the function is

$$\text{min}(x, na.rm = FALSE)$$

where *x* is a vector of numeric or logical values.

```
x <- c(30, NA, 12, 100, 43, 67, 100)
min(x, na.rm = FALSE)
```

```
[1] NA
```

```
min(x, na.rm = TRUE)
```

```
[1] 12
```

The minimum can be taken over more than one vector. To do so, list all desired vectors in the *min()* function, separating them with commas.

```
y <- c(0, 1000, 15, 20, 11, 0, 0)
min(x, y, na.rm = TRUE)
```

```
[1] 0
```

The minimum of a logical vector can be found. Using the usual interpretation of *TRUE* and *FALSE*, if the minimum of a vector equals 0, this implies that the event never occurred.

```
logical.x <- c(TRUE, TRUE, FALSE, TRUE)
min(logical.x)
```

```
[1] 0
```

Parallel Minimum: The *pmin()* function is used to determine pairwise minimum over a pair of vectors. *pmin()* determines the corresponding minimums. The basic structure for the pairwise maximum is

$$\text{pmax}(x, y, na.rm = FALSE)$$

where *x, y* are the given vector of numeric or logical values. The third argument *na.rm =* determines whether missing values should be ignored or not. By default, they are not ignored in the computation.

```
x <- c(30, NA, 12, 100, 43, 67, 100)
y <- c(0, 1000, 15, 20, 11, 0, 0)
pmin(x, y, na.rm = FALSE)
```

```
[1] 0 NA 12 20 11 0 0
```

```
pmin(x, y, na.rm = TRUE)
```

```
[1] 0 1000 12 20 11 0 0
```

Index of Minimum: The *which.min()* determines the index of the *first* occurrence of the minimum of a vector.

```
which.min(x)
```

```
[1] 4
```

Rounding

R has several methods for rounding output. Most of the methods given here change the computed value as opposed to just changing the displayed value.

The `round()` function is the basic rounding function. Beyond the given vector x of values to be rounded, its other argument indicates the number of digits. When rounding a number that ends in a 5, it may not round exactly as you expect. The operating system of the computer and the exact representation of the number can affect the result. Depending on the pairing, `round(0.15, digits = 1)` could be 0.1 or 0.2.

```
x <- c(-1.24, -2.75, 1.24, 2.54, 0.15)
round(x, digits = 1)

[1] -1.2 -2.8 1.2 2.5 0.1
```

The `floor()` and `ceiling()` functions round to an integer value. `floor()` rounds to the largest integer less than the given value (towards negative infinity). `ceiling()` rounds to the smallest integer greater than the given value (towards positive infinity).

```
floor(x)

[1] -2 -3 1 2 0

ceiling(x)

[1] -1 -2 2 3 1
```

The `trunc()` function returns an integer by dropping all decimal places.

```
trunc(x)

[1] -1 -2 1 2 0
```

Ranking & Sorting

Sorting:

Given a vector whose elements can be ordered, the `sort()` function can be used to rearrange the vector from 'smallest' to 'largest', or the other way around. If the results need to be used subsequently, the result needs to be assigned to a variable.

```
x <- c(30, 20, 50, NA, 10, 40)
sort(x, decreasing = FALSE, na.last = NA)

[1] 10 20 30 40 50
```

`decreasing` is a logical argument. When set to `FALSE`, the values are rearranged in increasing order. `na.last` determines what becomes of `NA` values. When set to `NA`, they are removed. Resulting in a shorter vector than one started with. When set to `TRUE` they are put at the end of the returned vector. When set to `FALSE`, they go to the beginning.

Ordering by Index:

The `order()` function does not sort a vector. However, its result can be used to sort a given vector. `order()` returns a rearrangement of the indexes (a permutation) of the given vector x that can be used to sort the given vector.

Suppose the first value in the given vector x is the fifth smallest value in x . The result from `order()` will have a 1 in the fifth position. Next, suppose the second value in x is the smallest value in x . `order` will place a 2 in the first position of its result.

```
order(x)

[1] 5 2 1 6 3 4

x[order(x)]

[1] 10 20 30 40 50 NA
```

This is useful when it is desired to sort a dataframe based upon one variable, or a few variables. In this first case, the result of `order()` can be used to rearrange the rows in the dataframe according to the values in `dat$V2` alone.

```
dat <- data.frame(V1 = LETTERS[1:6], V2 = c(7, 1, 6, 4, 1, 7), V3 = c(F, T, T, T, F, T))
dat

  V1 V2  V3
1  A  7 FALSE
2  B  1  TRUE
3  C  6  TRUE
4  D  4  TRUE
5  E  1 FALSE
6  F  7  TRUE

order.1.var <- order(dat$V2)
order.1.var

[1] 2 5 4 3 1 6

dat

  V1 V2  V3
1  A  7 FALSE
2  B  1  TRUE
3  C  6  TRUE
4  D  4  TRUE
5  E  1 FALSE
6  F  7  TRUE

LETTERS[1:6]

[1] "A" "B" "C" "D" "E" "F"

V2 = c(7, 1, 6, 4, 1, 7)
```

Notice that the tied rows are taken in the order they appear in the dataframe. However, if a second vector `datV3` is added to `order()`, it is used to adjust the output so that the tied rows have an order. Additional vectors can be added if needed. The `order()` function has an *na.last* and *decreasing*.

```
order.2.var <- order(dat$V2, dat$V3)
order.2.var

[1] 5 2 4 3 1 6

dat[order.2.var, ]

  V1 V2  V3
5  E  1 FALSE
2  B  1  TRUE
4  D  4  TRUE
3  C  6  TRUE
1  A  7 FALSE
6  F  7  TRUE
```

Ranking:

The result of the `rank()` function can't be used to sort a vector, or the rows of a dataframe. Instead, it provides the position of each value in each given value, if the values were sorted. In other words, it tells you the placement of each value in an ordered list.

```
x <- c(30, 20, 50, NA, 10, 40, 40)
sort(x)

[1] 10 20 30 40 40 50

rank(x, na.last = TRUE, ties.method = "average")

[1] 3.0 2.0 6.0 7.0 1.0 4.5 4.5
```

The argument *ties.method* determines how to rank ties. The default is to average them. In this example, forty is the fourth and fifth value in the ordered list. This averages to 4.5. Other options are “first”, “last”, “random”, “max”, and “min”. The first two put tied ranks into increasing or decreasing order, respectively. “max” uses the maximum rank for each set of tied values. While “min” uses the minimum.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

Questions. Questions. Questions. I need more questions

Chapter 25: Functions - Probability Distributions

R contains several functions pertaining to probability distributions for many types on random variables X . These functions relate to discrete random variables with finite and infinite sample spaces, and continuous random variables. Most of the functions discussed here come in sets of four that follow a specific naming pattern. Each function in a set starts with either a *d*, *p*, *q*, or *r*, followed by an abbreviation for the distributions name. (*xxx* is a placeholder for the abbreviation.) Each function has arguments that are specific to the distribution, and arguments that are common.

- dxxx** - density or mass function with common argument (x). This function determines the value of a density, or mass, function at the values x .
- pxxx** - distribution function with common arguments (q , *lower.tail*, *log.p*). This function determines cumulative probability $P(X \leq q)$ at q when the *lower.tail* is set to *TRUE*. When, the *lower.tail* argument is set to *FALSE*, the complementary probabilities $P(X > q)$ at q are determined. *log.p* controls whether a probability is returned (*FALSE*) or its log is returned (*TRUE*).
- qxxx** - quantile function with common arguments (p , *lower.tail*, *log.p*). For each value in p , this function returns a quantile q such that the cumulative probability at p is q when *lower.tail* is set to *TRUE*. In other words, this function solves $P(X \leq q) = p$ for q . When *lower.tail* is set to *FALSE*, the complement of the cumulative probability is used. The equation $P(X > q) = p$ is solved for q . *log.p* indicates whether a probability p is given (*FALSE*) or if its log is given (*TRUE*).
- rxxx** - random number generator with common argument (n). The random number generator produces a sample of n observations from the indicated probability distribution.

Regardless of the distribution, each type of function has some common elements. The first element q for the distribution function is a vector of quantiles.

Continuous Distributions

Several different continuous distributions have functions in *R*.

For the **normal** distribution, the functional names are *dnorm*(), *pnorm*(), *qnorm*(), and *rnorm*(). The parameters for a normal distribution are indicated by the *mean* and *sd* (standard deviation) arguments. The default is a standard normal distribution.

```
dnorm(x, mean = 0, sd = 1, log = FALSE)
pnorm(q, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)
qnorm(p, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)
rnorm(n, mean = 0, sd = 1)
```

For the **chi-square** distribution, the functional names are *dchisq*(), *pchisq*(), *qchisq*(), and *rchisq*(). The parameters are indicated by the *df* (degrees of freedom) and *ncp* (non-centrality parameter) arguments. *ncp* defaults to 0.

```
dchisq(x, df, ncp = 0, log = FALSE)
pchisq(q, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
qchisq(p, df, ncp = 0, lower.tail = TRUE, log.p = FALSE)
rchisq(n, df, ncp = 0)
```

For the **Student's t** distribution, the functional names are *dt*(), *pt*(), *qt*(), and *rt*(). The parameters are indicated by the *df* (degrees of freedom) and *ncp* (non-centrality parameter) arguments. The *ncp* argument was omitted to indicate that the centralized *t* distribution is used.

```
dt(x, df, log = FALSE)
pt(q, df, lower.tail = TRUE, log.p = FALSE)
qt(p, df, lower.tail = TRUE, log.p = FALSE)
rt(n, df)
```

For the **uniform** distribution, the functional names are *dunif*(), *punif*(), *qunif*(), and *runif*(). The parameters are indicated by the *min* (minimum) and *max* (maximum) arguments. As defaults, these are set to zero and one, respectively.

```
dunif(x, min = 0, max = 1, log = FALSE)
punif(q, min = 0, max = 1, lower.tail = TRUE, log.p = FALSE)
qunif(p, min = 0, max = 1, lower.tail = TRUE, log.p = FALSE)
runif(n, min = 0, max = 1)
```

Discrete Distributions

For the **binomial** distribution, the functional names are `dbinom()`, `pbinom()`, `qbinom()`, and `rbinom()`. The parameters indicated by the *size* (number of trials) and *prob* (probability of success on an individual trial) arguments.

```
dbinom(x, size, prob, log = FALSE)
pbinom(q, size, prob, lower.tail = TRUE, log.p = FALSE)
qbinom(p, size, prob, lower.tail = TRUE, log.p = FALSE)
rbinom(n, size, prob)
```

For the **poisson** distribution, the functional names are `dpois()`, `ppois()`, `qpois()`, and `rpois()`. The sole parameter is indicated by the *lambda* (mean) argument.

```
dpois(x, lambda, log = FALSE)
ppois(q, lambda, lower.tail = TRUE, log.p = FALSE)
qpois(p, lambda, lower.tail = TRUE, log.p = FALSE)
rpois(n, lambda)
```

In the help menu, searching for “distributions” finds a listing of many more distributions that are included included in *base R*.

Sample Function

While the random number generators in the previous two functions will generate values from two ‘named’ distributions, the `sample()` function can be used to generate values from any discrete random variable with a finite sample space. To do so, a few arguments must be set. *x* indicate all the possible values in the sample space. *prob* indicates the corresponding probabilities for each value in *x*. *size* indicates the number of elements of *x* to be sampled. *replace* is a logical argument. Setting *replace* to *TRUE* indicates that values taken from *x* should be replaced after each selection. This will simulate an independent and identically distributed sample from *x*. Setting *replace* to *FALSE* indicates that each value in *x* will not be allowed to be selected more than once. Effectively, when *replace* = *FALSE*, *size* can not be any larger than the number of elements in *x*.

```
possible.values <- c(5, 7, 69, 80)
related.probabilities <- c(0.1, 0.2, 0.3, 0.4)
sample.size <- 10

sample(x = possible.values, size = sample.size, replace = TRUE, prob = related.probabilities)

[1] 80 7 80 5 80 69 80 80 7 80
```

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.

Questions. Questions. Questions. I need more questions

Chapter 26: Basic Statistics Functions

Sample Mean

The `mean()` function computes the average of all values in a vector³¹. `mean()` is an example of a function with an `na.rm` logical argument. Its default setting is `FALSE`. This can lead to an `NA` result, if there are missing values in the vector being averaged.

```
x <- c(1, 5, 5, 5, 5, NA, 5, 5, 5, 9)
mean(x)

[1] NA

mean(x, na.rm = TRUE)

[1] 5
```

Sample Variance & Sample Standard Deviation

The `var()` computes the (unbiased) **sample variance** of a given vector. This is the sample variance whose denominator is one less than the number of data values³². Corresponding to this is the `sd()` function that computes the **sample standard deviation**. As with `mean()`, missing values `NA` can affect the result. As such, `na.rm` is an argument for these two functions.

```
var(x, na.rm = TRUE)

[1] 4

sd(x, na.rm = TRUE)

[1] 2
```

Sample Quantiles

Sample quantiles are computed using `quantile()`. Its first argument is a vector of values. The `probs` argument is a numerical vector identifying the quantiles to be computed. The default is to produce a five-number summary. `na.rm` is available to handle missing values. The `type` argument is used to determine the method used for interpolating between consecutive order statistics.

```
x <- c(1:10, NA)
quantile(x, probs = c(0, 0.2, 0.5, 0.76, 1), na.rm = TRUE, type = 6, names = TRUE)

 0%   20%   50%   76%  100%
1.00  2.20  5.50  8.36 10.00
```

Sample Range

The `range()` function produces a two element vector containing the largest and smallest values of its given arguments³³. The `na.rm` argument is available to handle missing values. The `finite` determines if infinite values should be ignored.

```
range(1:10, NA, Inf, na.rm = TRUE) #multiple arguments are given

[1] 1 Inf

range(1:10, NA, Inf, na.rm = TRUE, finite = TRUE)

[1] 1 10
```

Character data can also be used.

```
range(letters, NA, Inf, na.rm = TRUE) #multiple arguments are given

[1] "a" "z"
```

³¹ And many other R objects

³² If n is the length of the vector, the denominator is $n - 1$.

³³ vectors, dataframes, matrices, etc.

Frequencies

The `table()` will produce a contingency table for the elements of the objects³⁴ given to it. With a single argument, it produces a one-way table of frequencies

```
x <- c(1, 2, 4, 3, 2, 3, 4, 3, 2, 1, 2, 3)
table(x)

x
1 2 3 4
2 4 4 2
```

When two arguments are given, a two-way table of frequencies is produced.

```
y <- c("a", "b", "b", "a", "c", "b", "a", "b", "a", "c", "b", "b")
table(x, y)

      y
x     a b c
1  1 0 1
2  1 2 1
3  1 3 0
4  1 1 0
```

It is important to look at the structure of the result. In this case, a **table** is returned. Values within this result can be retrieved by treating this result as an array.

```
str(table(x, y))

'table' int [1:4, 1:3] 1 1 1 1 0 2 3 1 1 1 ...
- attr(*, "dimnames")=List of 2
 ..$ x: chr [1:4] "1" "2" "3" "4"
 ..$ y: chr [1:3] "a" "b" "c"
```

Data Summaries

The `summary()` function is generic function whose results will depend on the type of *R* object is given. This example makes use of a dataframe.

```
numbers <- 1:10
logicals <- c(T, T, T, T, F, F, T, F, F, T)
dat <- data.frame(numbers, logicals)
summary(dat)

      numbers      logicals
Min.   : 1.00   Mode :logical
1st Qu.: 3.25  FALSE:4
Median : 5.50   TRUE :6
Mean    : 5.50
3rd Qu.: 7.75
Max.    :10.00
```

This produced the mean and the five-number summary of the numerical data, and the frequencies of the logical data.

³⁴ Any objects that can viewed as factors.

Exercises

The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.

Unless directed otherwise, a result should always be displayed.


Questions. Questions. Questions. I need more questions

Chapter 27: Patchwork

patchwork is a package used for combining graphs into a single image and controlling their layout. *patchwork* is mainly used in conjunction with graphs produced by a package called *ggplot2*³⁵.

```
library(patchwork)
```

When graphs are to be laid out using *patchwork*, they will have already been created using *ggplot2* and each graph will have had a variable name assigned to it. For the following examples, these four graphs will be used. Their variable name is noted in the upper left hand corner of each graph³⁶



Basic Operations

A great deal of control over the layout of multiple graphs comes from three basic operators: “|”, “/” and “()”. The “|” places


The “|” sign place graphs side-by-side. The “|” was used to make the first graph on this page. This example places the four graphs in a row. Each is scaled down.

```
graph1 | graph2 | graph3 | graph4 | graph1 | graph2 | graph3 | graph4
```



The “/” sign is used to stack graphs on top of one another.

```
graph1/graph2
```




The parenthesis are used to group graphs together when the “|” and “/” commands are used together. With one grouping, *graph1* and *graph2* are stacked. Their stacked image is placed next to *graph3*.

³⁵ The *ggplot2* package will be covered in the next chapter.


³⁶ For consistency, each set of graphs will be displayed in a region that is the same height and width as this first graph. This may make some graphs look odd. However, the layout is the current concern.

```
(graph1/graph2) | graph3
```



While another grouping places *graph1* above a side-by-side pairing of *graph2* and *graph3*.


```
graph1/(graph2 | graph3)
```



Other Operations

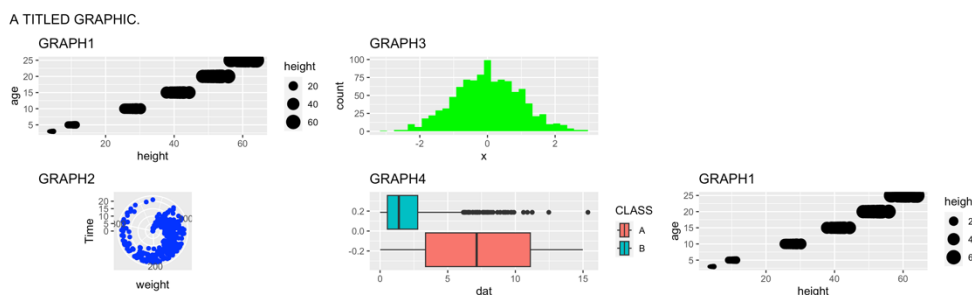
Other operations can be used. The “+” can be used to make a grid of graphics. When used in conjunction with the `plot_layout()` function the dimension and fill of the grid can be controlled³⁷.

```
graph1 + graph2 + graph3 + graph4 + graph1 + plot_layout(nrow = 2, ncol = 3, byrow = FALSE)
```



The `plot_spacer()` command can be used to skip a graphing position. While adding a `widths` argument to `plot_layout` will even out the widths of the graphs, if possible. `plot_annotation()` can add an overall title.

```
graph1 + graph2 + graph3 + graph4 + plot_spacer() + graph1 + plot_layout(nrow = 2, ncol = 3, byrow = FALSE, widths = 1) + plot_annotation(title = "A TITLED GRAPHIC.")
```



³⁷ This is much the same as the `matrix()` function.

Exercises


The exercises given below should be completed using only information given in the current section and previous sections. The **help** menu should not be used as a window into all possible functions unless it is specifically allowed. Similarly, only those packages that have been introduced in the current or previous sections may be used. Any computations that need to be completed should be done making use of R and not in your head. For example, if asked to find the sum of one and two, you should type in `1+2`, and not simply give the answer of 3.

When told to use a specific number of lines of code, this should be interpreted as one line is one complete command. Some commands can be very long and necessary will wrap around. In some instances, a command is purposely wrapped around several lines. This insures readability.


Unless directed otherwise, a result should always be displayed.

Recreate each of these graphics using the graphs named: `patchwork01`, `patchwork02`, `patchwork03`, `patchwork04`.


1. Graphic #01



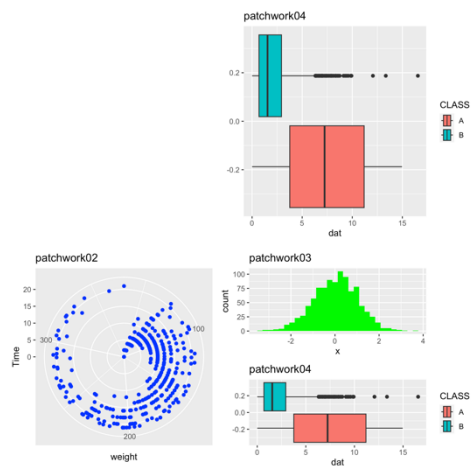
2. Graphic #02




3. Graphic #03



4. Graphic #04

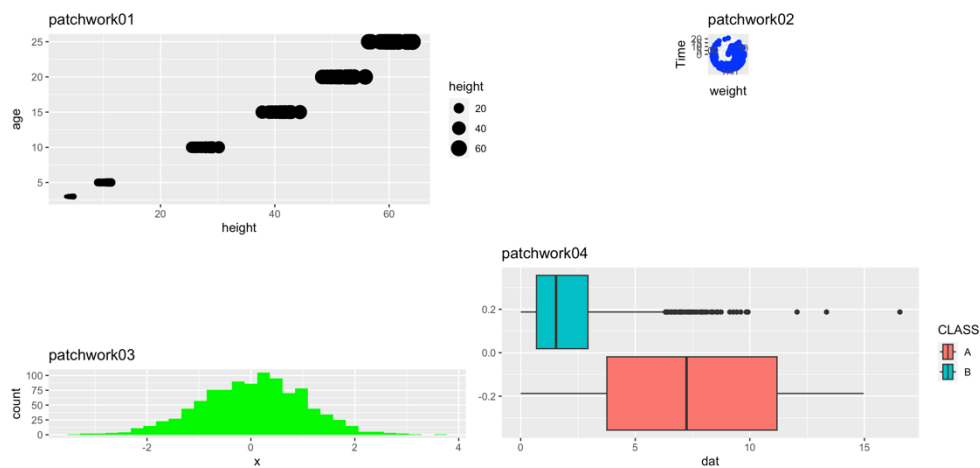


5. Graphic #05



6. Graphic #06

The hard one.



Chapter 28: Bibliography (Not Finished Yet)